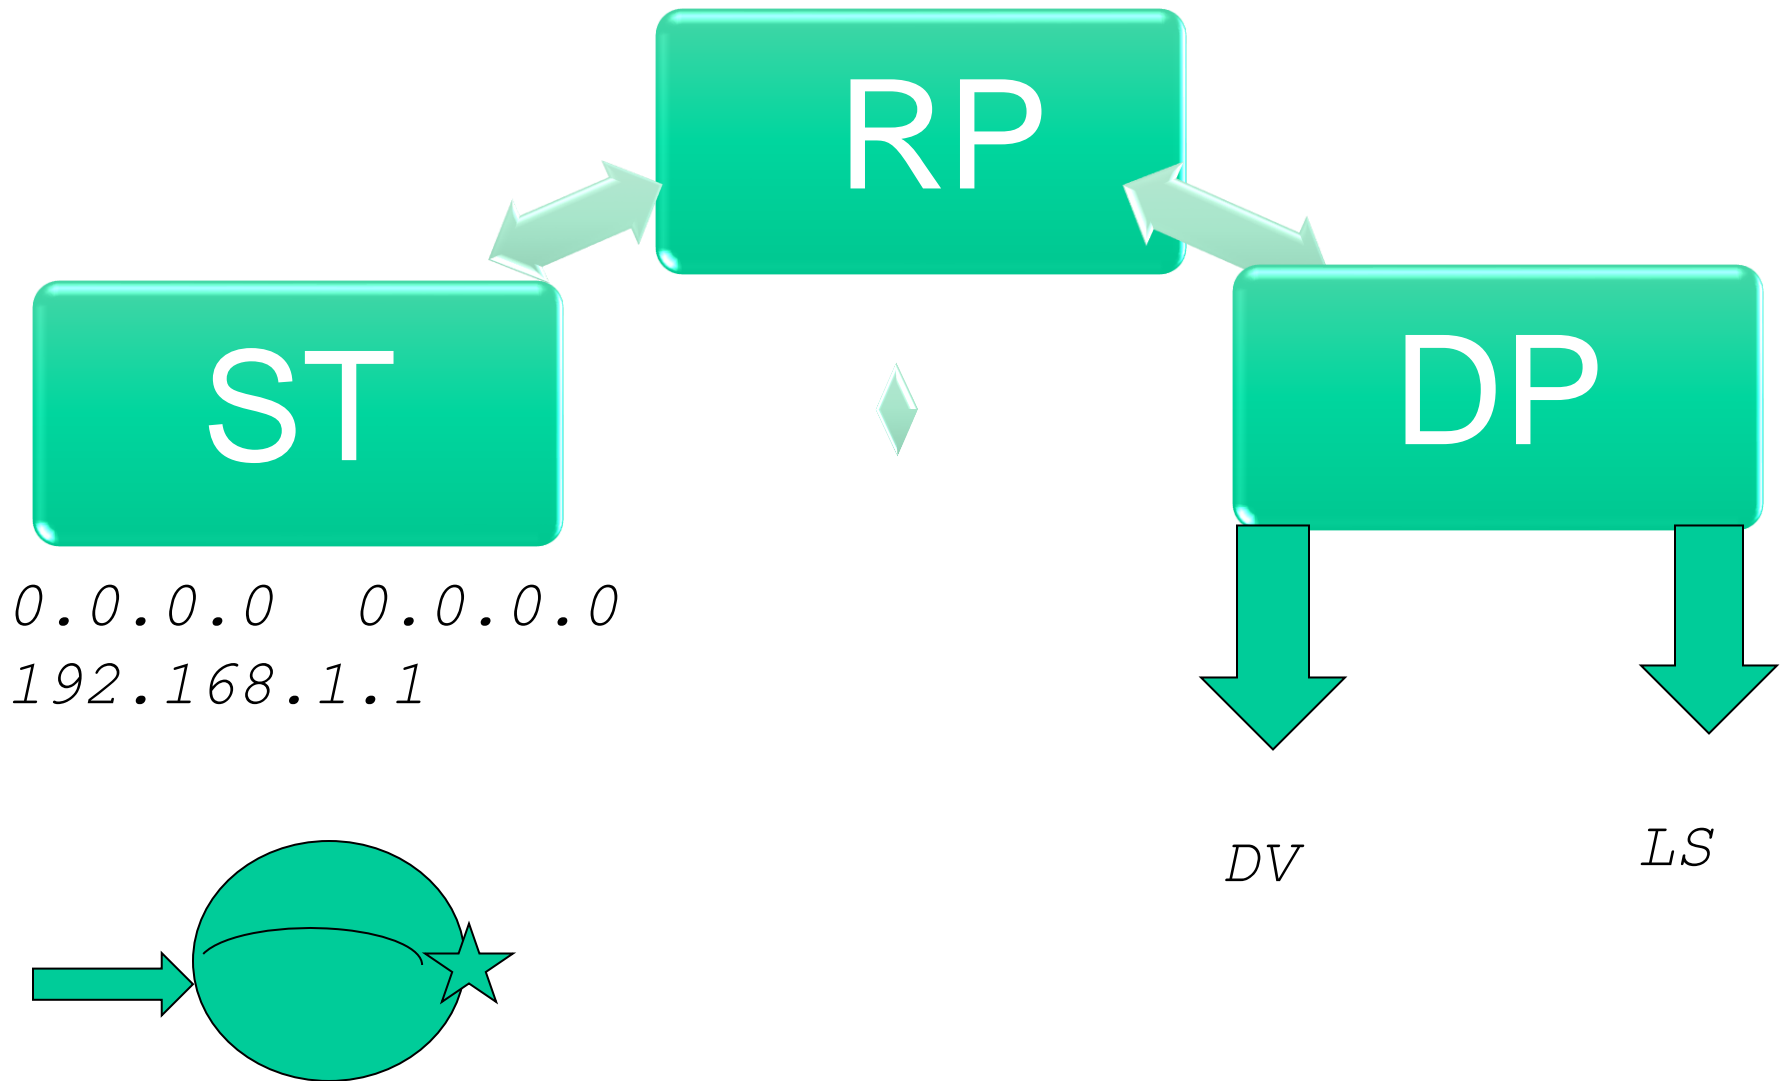



# Dynamic Routing Protocols I

## RIP

---

The first module on dynamic routing protocols. This module provides an overview of routing, introduces terminology (interdomain, intradomain, autonomous system),



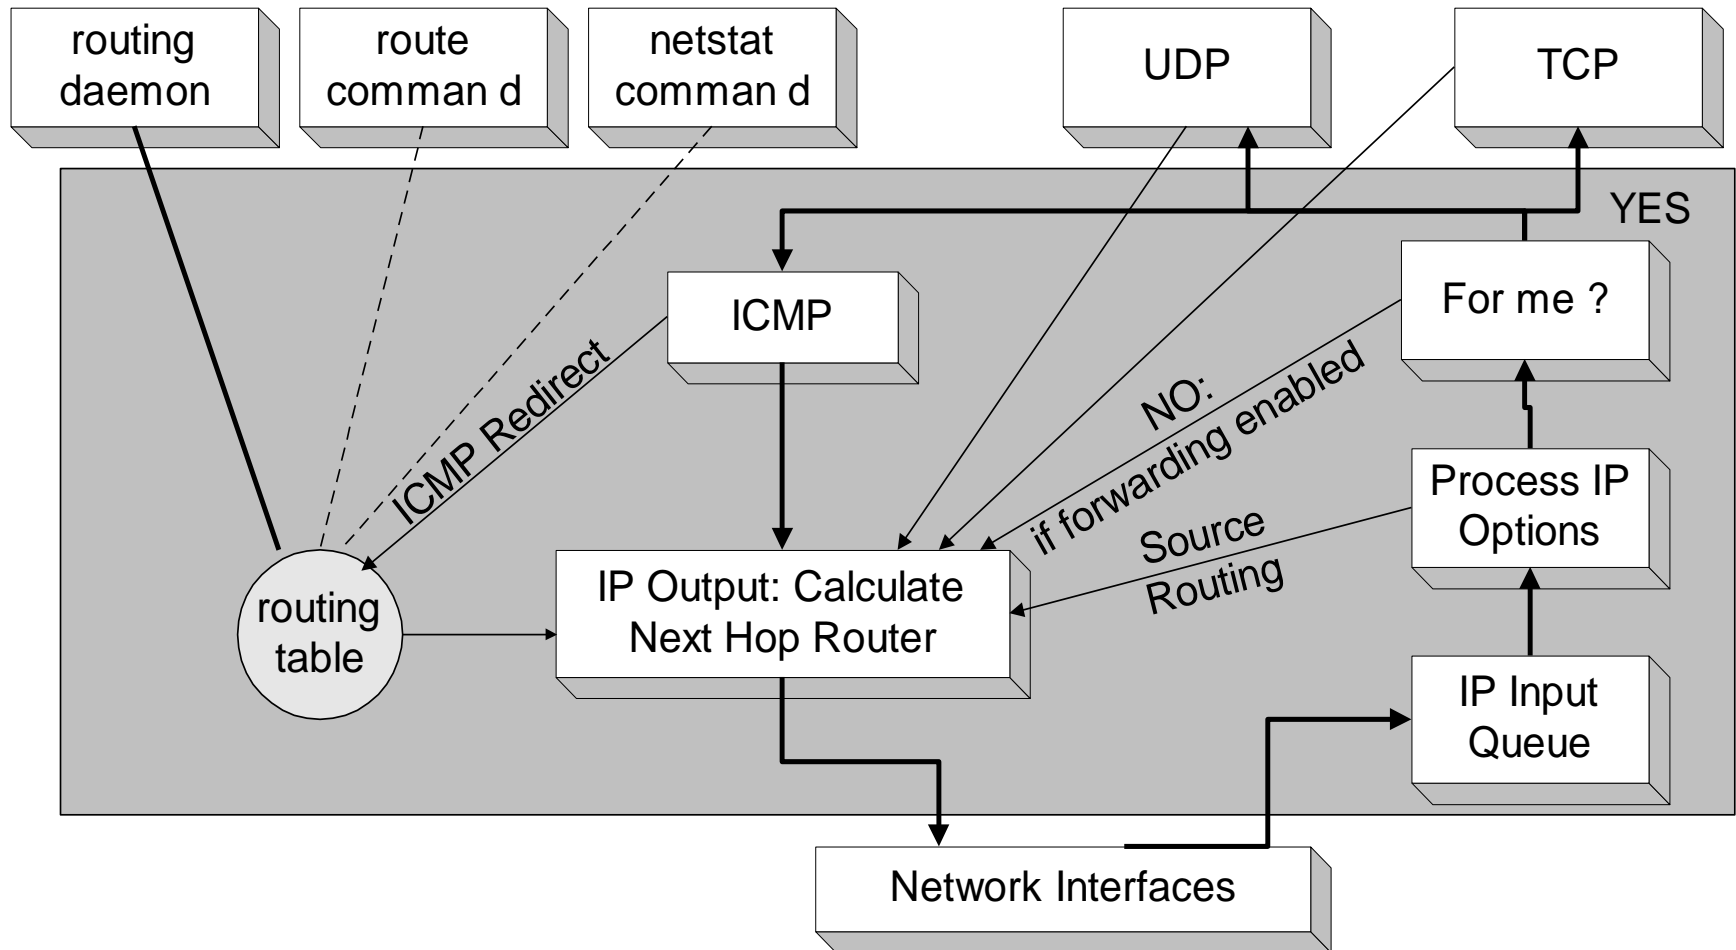


$$F_{X \rightarrow Y} = AB + CD + AED + CEB$$

# Routing

- **Recall:** There are two parts to routing IP packets:
  1. How to pass a packet from an input interface to the output interface of a router (packet forwarding) ?
  2. How to find and setup a route ?
- We already discussed the packet forwarding part
- There are two approaches for calculating the routing tables:
  - Static Routing
  - Dynamic Routing: Routes are calculated by a routing protocol

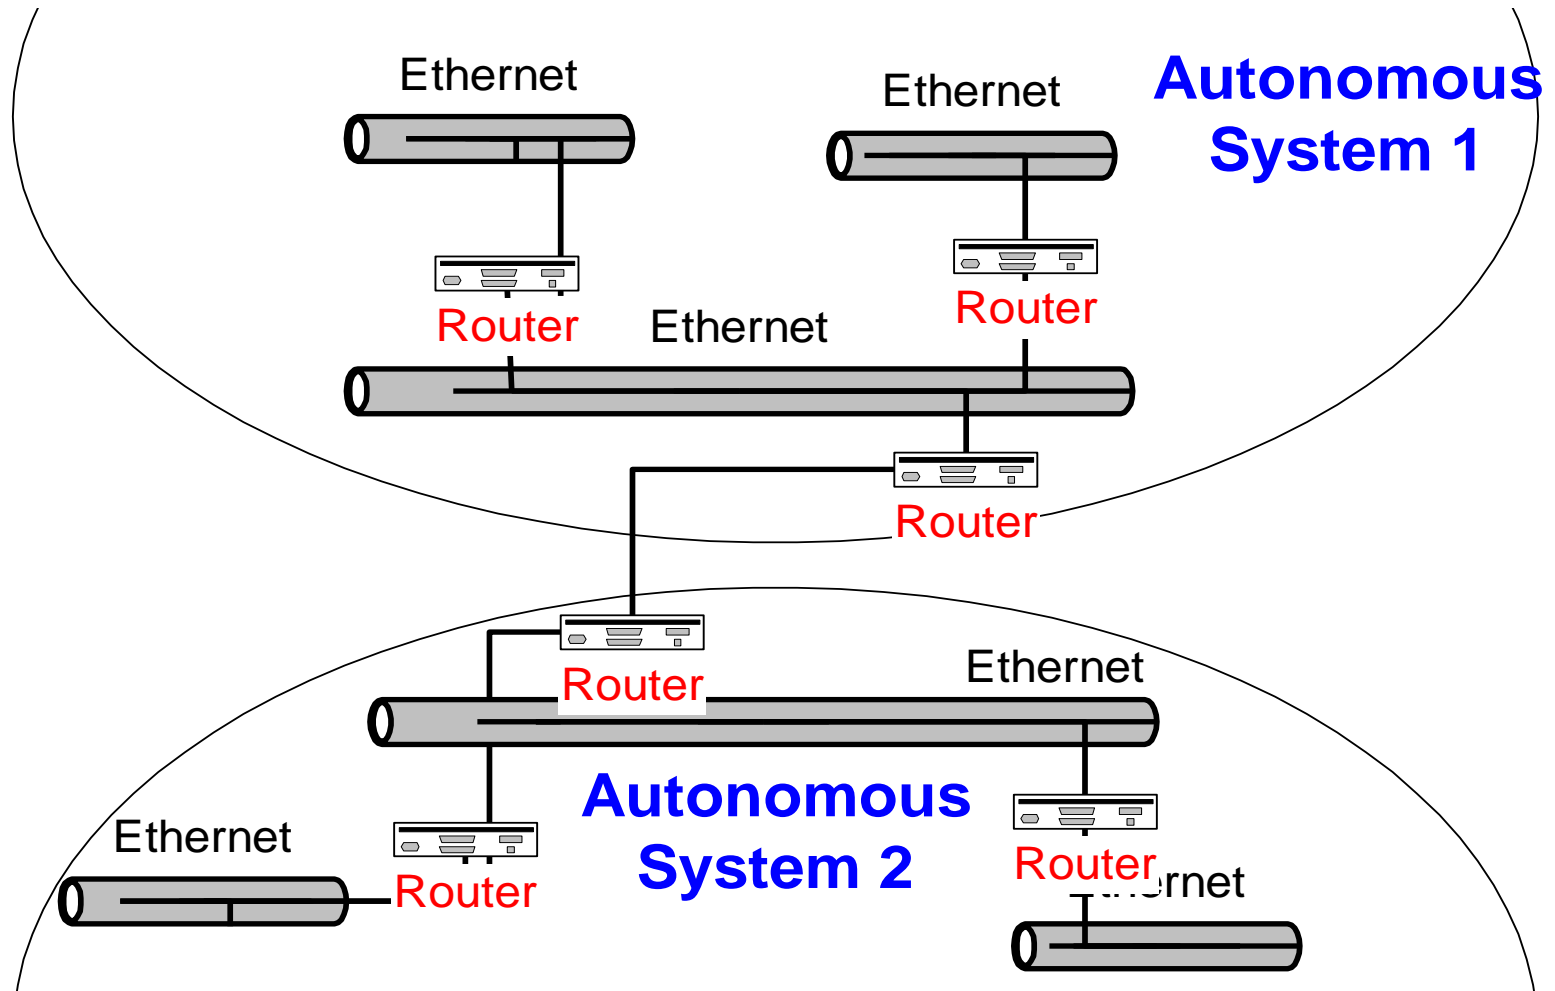
# IP Routing




# Autonomous Systems

- An **autonomous system** is a region of the Internet that is administered by a single entity.
- Examples of autonomous regions are:
  - UVA's campus network
  - MCI's backbone network
  - Regional Internet Service Provider
- Routing is done differently within an autonomous system (**intradomain routing**) and between autonomous system (**interdomain routing**).

# Autonomous Systems (AS)



# Interdomain and Intradomain Routing



## Intradomain Routing

- **Routing within an AS**
- Ignores the Internet outside the AS
- Protocols for Intradomain routing are also called **Interior Gateway Protocols** or **IGP's**.
- Popular protocols are
  - RIP (simple, old)
  - OSPF (better)

## Interdomain Routing

- **Routing between AS's**
- Assumes that the Internet consists of a collection of interconnected AS's
- Normally, there is one dedicated router in each AS that handles interdomain traffic.
- Protocols for interdomain routing are also called **Exterior Gateway Protocols** or **EGP's**.
- Routing protocols:
  - EGP
  - BGP (more recent)



# Components of a Routing Algorithm

- A procedure for sending and receiving reachability information about network to other routers
- A procedure for calculating optimal routes
  - Routes are calculated using a shortest path algorithm:
    - **Goal: Given a network where each link is assigned a cost. Find the path with the least cost between two networks with minimum cost.**
- A procedure for reacting to and advertising topology changes

# Approaches to Shortest Path Routing

- There are two basic routing algorithms found on the Internet.

## 1. Distance Vector Routing

- Each node **knows the distance (=cost) to its directly connected neighbors**
- A node sends periodically a list of routing updates to its neighbors.
- If all nodes update their distances, the routing tables eventually converge
- New nodes advertise themselves to their neighbors

## 2. Link State Routing

- **Each node knows the distance to its neighbors**
- The distance information (**=link state**) is broadcast to all nodes in the network
- **Each node calculates the routing tables independently**

# Routing Algorithms in the Internet

## Distance Vector

- **Routing Information Protocol (RIP)**
- Gateway-to-Gateway Protocol (GGP)
- Exterior Gateway Protocol (EGP)
- Interior Gateway Routing Protocol (IGRP)

***Bellman  
algorithm***

## Link State

- Intermediate System - Intermediate System (IS-IS)
- **Open Shortest Path First (OSPF)**

***Dijkstra  
Algorithm***

# Dynamic IP Routing Protocols


- In Unix systems, the dynamic setting of routing tables is done by the **routed** or **gated** daemons
- The routing daemons execute the following intradomain and interdomain routing protocols

<i>Daemon</i>	<i>Hello</i>	<i>RIP</i>	<i>OSPF</i>	<i>EGP</i>	<i>BGP</i>
<b>routed</b>		V1			
<b>Gated</b> (Version 3)	Yes	V1 V2	V2	Yes	V2, V3

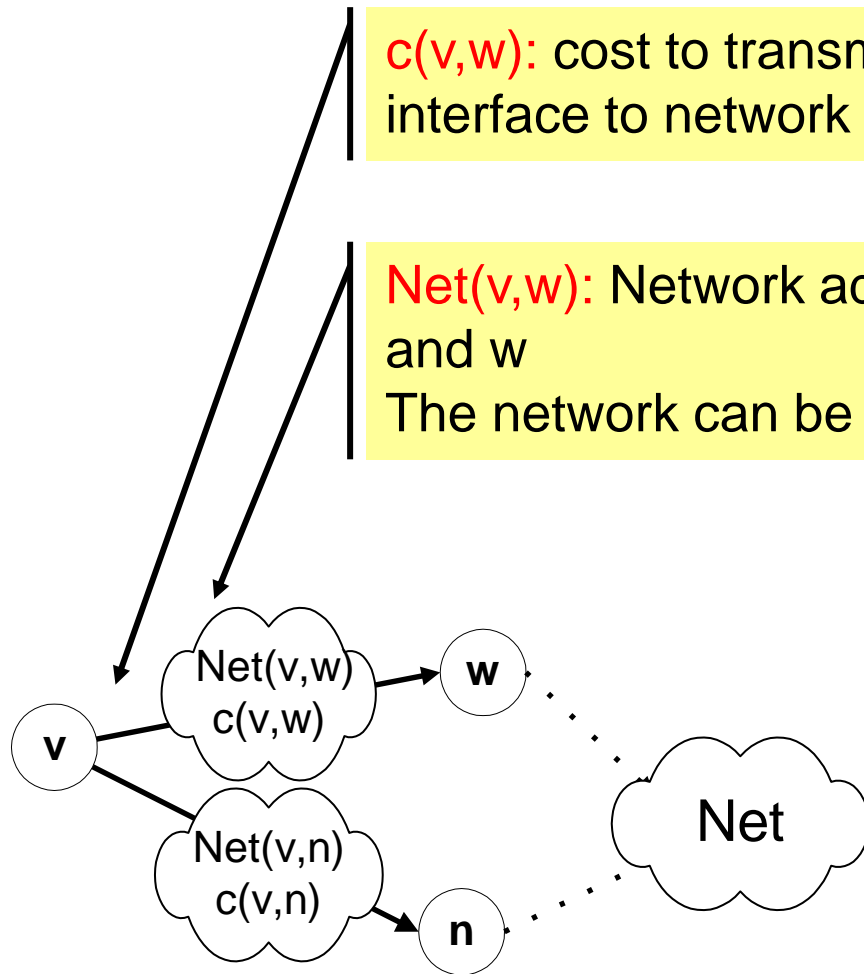
# A network as a graph

- In the following, networks are represented as a network graph:
  - nodes are connected by networks
    - network can be a link or a LAN
  - network interface has cost
  - networks are destinations
  - $\text{Net}(v,w)$  is an IP address of a network

- For ease of notation, we often replace the clouds between nodes by simple links.



# Distance Vector Algorithm: Routing Table



$c(v,w)$ : cost to transmit on the interface to network  $\text{Net}(v,w)$

$\text{Net}(v,w)$ : Network address of the network between  $v$  and  $w$   
The network can be a link, but could also be a LAN

**RoutingTable of node  $v$**


Dest	via (next hop)	cost
Net	$n$	$D(v, \text{Net})$

# Distance Vector Algorithm: Messages

RoutingTable of node v

Dest	via (next hop)	cost
Net	n	$D(v, \text{Net})$

- Nodes send messages to their neighbors which contain routing table entries




- A message has the format:  $[\text{Net} , D(v, \text{Net})]$  means “*My cost to go to Net is  $D(v, \text{Net})$* ”

# Distance Vector Algorithm: Sending Updates

## RoutingTable of node v

Dest	via (next hop)	cost
Net <sub>1</sub>	m	D(v, Net <sub>1</sub> )
Net <sub>2</sub>	n	D(v, Net <sub>2</sub> )
...	...	...
Net <sub>N</sub>	w	D(v, Net <sub>N</sub> )

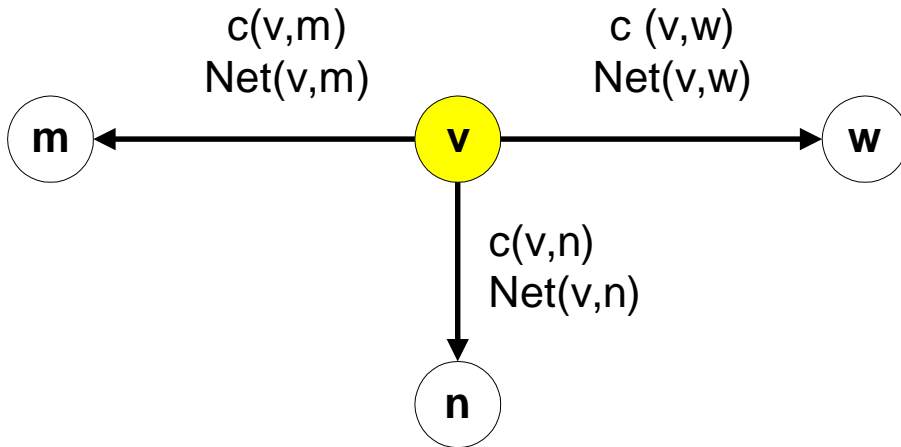
Periodically, each node  $v$  sends the content of its routing table to its neighbors:





# Initiating Routing Table I

- Suppose a new node  $v$  becomes active.
- The cost to access directly connected networks is zero:
  - $D(v, \text{Net}(v,m)) = 0$
  - $D(v, \text{Net}(v,w)) = 0$
  - $D(v, \text{Net}(v,n)) = 0$



**RoutingTable**


Dest	via (next hop)	cost
$\text{Net}(v,m)$	$m$	$0$
$\text{Net}(v,w)$	$w$	$0$
$\text{Net}(v,n)$	$n$	$0$

# Initiating Routing Table II

**RoutingTable**


Dest	via (next hop)	cost
Net(v,m)	m	0
Net(v,w)	w	0
Net(v,n)	n	0

- New node v sends the routing table entry to all its neighbors:




# Initiating Routing Table III

- Node  $v$  receives the routing tables from other nodes and builds up its routing table



# Updating Routing Tables I

- Suppose node  $v$  receives a message from node  $m$ :  $[\text{Net}, D(m, \text{Net})]$




Node  $v$  updates its routing table and sends out further messages if the message reduces the cost of a route:

```
if (  $D(m, \text{Net}) + c(v, m) < D(v, \text{Net})$  ) {  
     $D^{\text{new}}(v, \text{Net}) := D(m, \text{Net}) + c(v, m)$ ;  
    Update routing table;  
    send message  $[\text{Net}, D^{\text{new}}(v, \text{Net})]$  to all neighbors  
}
```

# Updating Routing Tables II

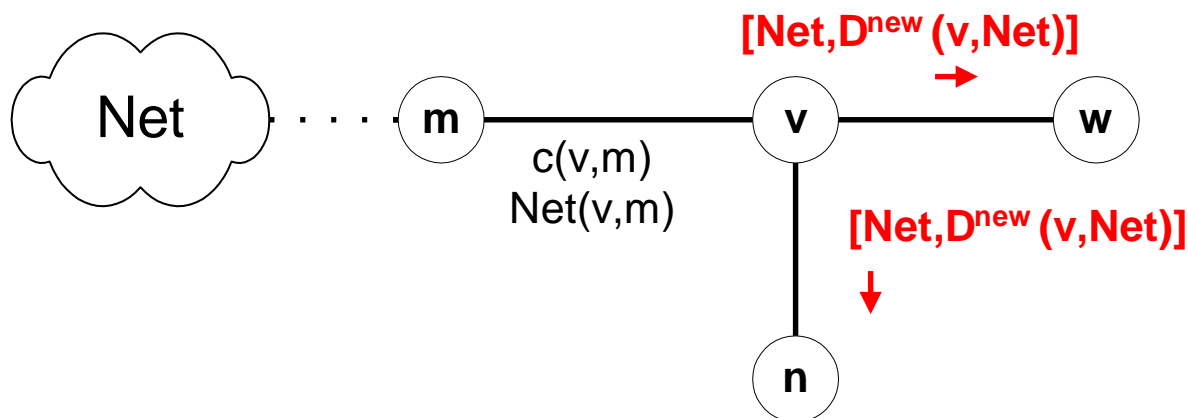
- Before receiving the message:



RoutingTable

Dest	via (next hop)	cost
Net	??	$D(v, \text{Net})$

- Suppose  $D(m, \text{Net}) + c(v, m) < D(v, \text{Net})$ :




RoutingTable

Dest	via (next hop)	cost
Net	m	$D^{\text{new}}(v, \text{Net})$


# Example

- Assume:
- link cost is 1, i.e.,  $c(v,w) = 1$
  - all updates, updates occur simultaneously
  - Initially, each router only knows the cost of connected interfaces



Net	via	cost		Net	via	cost		Net	via	cost		Net	via	cost
<b>t=0:</b>				<b>t=0:</b>				<b>t=0:</b>				<b>t=0:</b>		
10.0.1.0	-	0		10.0.2.0	-	0		10.0.3.0	-	0		10.0.4.0	-	0
10.0.2.0	-	0		10.0.3.0	-	0		10.0.4.0	-	0		10.0.5.0	-	0
			→ ←				→ ←				→ ←			
<b>t=1:</b>				<b>t=1:</b>				<b>t=1:</b>				<b>t=1:</b>		
10.0.1.0	-	0		10.0.1.0	10.0.2.1	1		10.0.2.0	10.0.3.1	1		10.0.3.0	10.0.4.1	1
10.0.2.0	-	0		10.0.2.0	-	0		10.0.3.0	-	0		10.0.4.0	-	0
10.0.3.0	10.0.2.2	1		10.0.3.0	-	0		10.0.4.0	-	0		10.0.5.0	-	0
			→ ←	10.0.4.0	10.0.3.2	1		10.0.5.0	10.0.4.2	1				
<b>t=2:</b>				<b>t=2:</b>				<b>t=2:</b>				<b>t=2:</b>		
10.0.1.0	-	0		10.0.1.0	10.0.2.1	1		10.0.1.0	10.0.3.1	2		10.0.2.0	10.0.4.1	2
10.0.2.0	-	0		10.0.2.0	-	0		10.0.2.0	10.0.3.1	1		10.0.3.0	10.0.4.1	1
10.0.3.0	10.0.2.2	1		10.0.3.0	-	0		10.0.3.0	-	0		10.0.4.0	-	0
10.0.4.0	10.0.2.2	2		10.0.4.0	10.0.3.2	1		10.0.4.0	-	0		10.0.5.0	-	0
			→ ←	10.0.5.0	10.0.3.2	2		10.0.5.0	10.0.4.2	1				

# Example



Net	via	cost	Net	via	cost	Net	via	cost	Net	via	cost
<b>t=2:</b>			<b>t=2:</b>			<b>t=2:</b>			<b>t=2:</b>		
10.0.1.0	-	0	10.0.1.0	10.0.2.1	1	10.0.1.0	10.0.3.1	2	10.0.2.0	10.0.4.1	2
10.0.2.0	-	0	10.0.2.0	-	0	10.0.2.0	10.0.3.1	1	10.0.3.0	10.0.4.1	1
10.0.3.0	10.0.2.2	1	10.0.3.0	-	0	10.0.3.0	-	0	10.0.4.0	-	0
10.0.4.0	10.0.2.2	2	10.0.4.0	10.0.3.2	1	10.0.4.0	-	0	10.0.5.0	-	0
			10.0.5.0	10.0.3.2	2	10.0.5.0	10.0.4.2	1			
<b>t=3:</b>			<b>t=3:</b>			<b>t=3:</b>			<b>t=3:</b>		
10.0.1.0	-	0	10.0.1.0	10.0.2.1	1	10.0.1.0	10.0.3.1	2	10.0.1.0	10.0.4.1	3
10.0.2.0	-	0	10.0.2.0	-	0	10.0.2.0	10.0.3.1	1	10.0.2.0	10.0.4.1	2
10.0.3.0	10.0.2.2	1	10.0.3.0	-	0	10.0.3.0	-	0	10.0.3.0	10.0.4.1	1
10.0.4.0	10.0.2.2	2	10.0.4.0	10.0.3.2	1	10.0.4.0	-	0	10.0.4.0	-	0
10.0.5.0	10.0.2.2	3	10.0.5.0	10.0.3.2	2	10.0.5.0	10.0.4.2	1	10.0.5.0	-	0


Now, routing tables have converged !

# Characteristics of Distance Vector Routing

- **Periodic Updates:** Updates to the routing tables are sent at the end of a certain time period. A typical value is 90 seconds.
- **Triggered Updates:** If a metric changes on a link, a router immediately sends out an update without waiting for the end of the update period.
- **Full Routing Table Update:** Most distance vector routing protocols send their neighbors the entire routing table (not only entries which change).
- **Route invalidation timers:** Routing table entries are invalid if they are not refreshed. A typical value is to invalidate an entry if no update is received after 3-6 update periods.



# The Count-to-Infinity Problem



A's Routing Table

to	via (next hop)	cost
C	B	2

B's Routing Table

to	via (next hop)	cost
C	C	1

now link B-C goes down

C	B	2
---	---	---

C	-	∞
---	---	---

C	2
---	---

C	∞
---	---

C	-	∞
---	---	---

C	A	3
---	---	---

C	∞
---	---

C	3
---	---

C	B	4
---	---	---

C	-	∞
---	---	---

C	4
---	---

C	∞
---	---

# Count-to-Infinity

---

- The reason for the count-to-infinity problem is that each node only has a “next-hop-view”
- For example, in the first step, A did not realize that its route (with cost 2) to C went through node B
- How can the Count-to-Infinity problem be solved?

# Count-to-Infinity

- The reason for the count-to-infinity problem is that each node only has a “next-hop-view”
- For example, in the first step, A did not realize that its route (with cost 2) to C went through node B
- How can the Count-to-Infinity problem be solved?
- **Solution 1:** Always advertise the entire path in an update message (**Path vectors**)
  - If routing tables are large, the routing messages require substantial bandwidth
  - BGP uses this solution

# Count-to-Infinity

- The reason for the count-to-infinity problem is that each node only has a “next-hop-view”
- For example, in the first step, A did not realize that its route (with cost 2) to C went through node B
- How can the Count-to-Infinity problem be solved?
- **Solution 2:** Never advertise the cost to a neighbor if this neighbor is the next hop on the current path (**Split Horizon**)
  - Example: A would not send the first routing update to B, since B is the next hop on A’s current route to C
  - Split Horizon does not solve count-to-infinity in all cases!

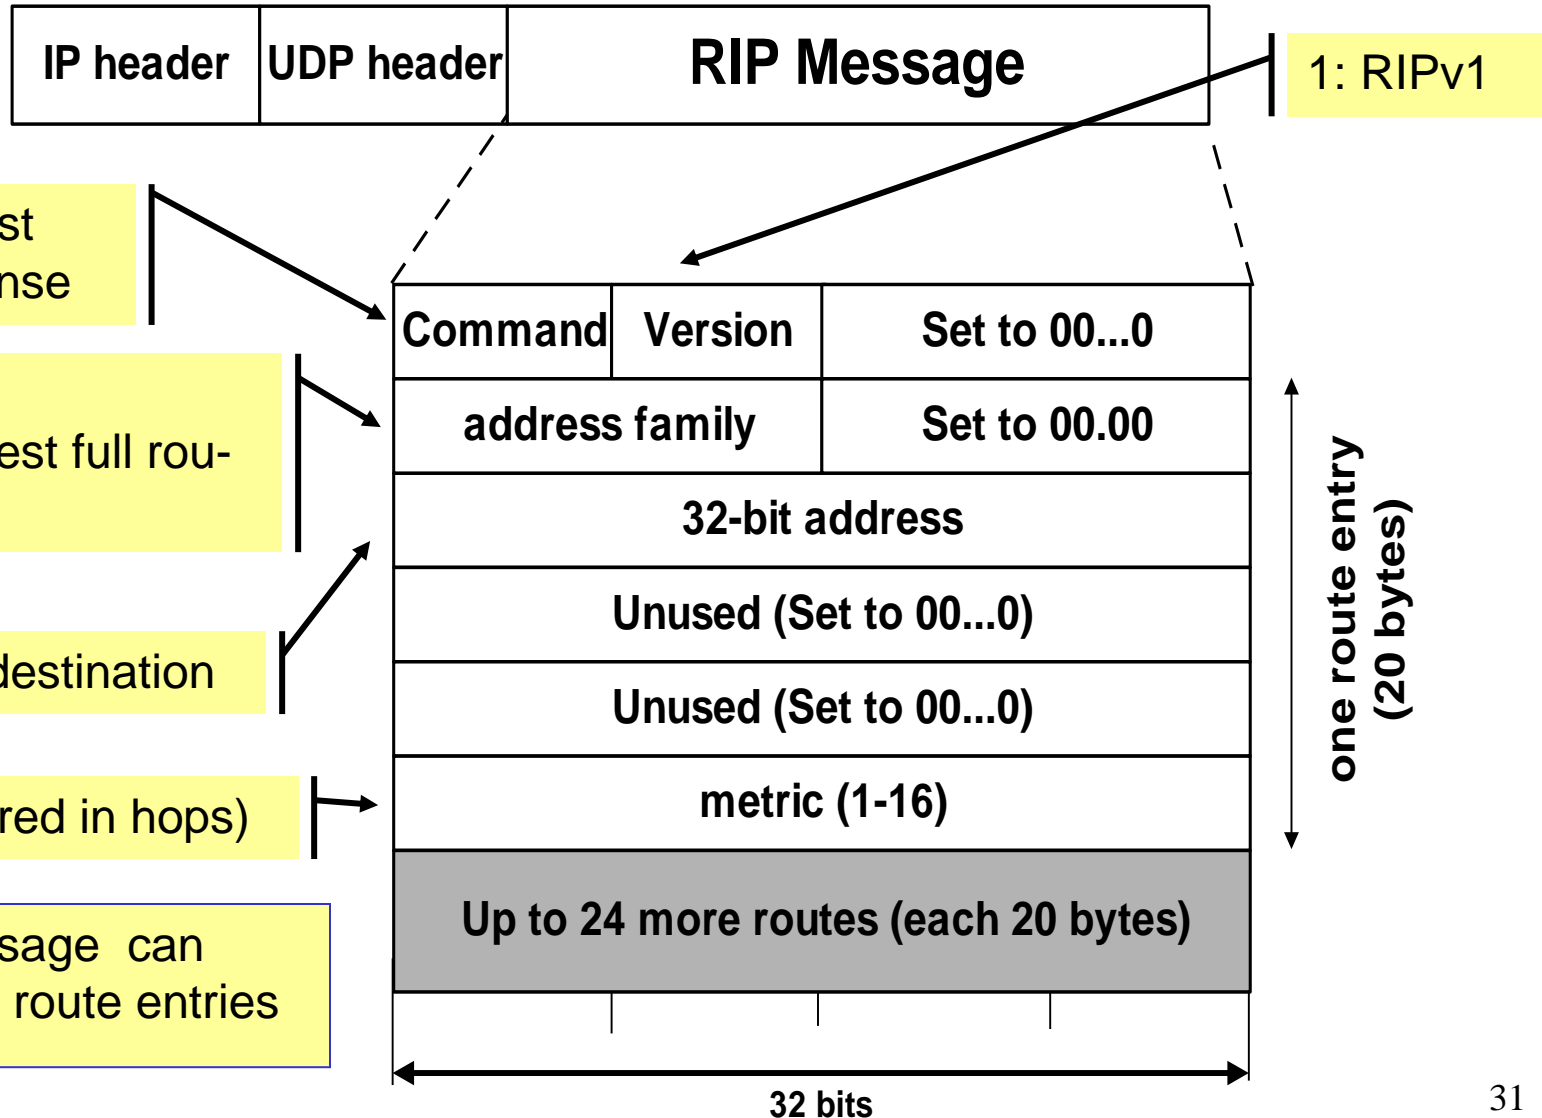
# RIP - Routing Information Protocol

- A simple intradomain protocol
- Straightforward implementation of Distance Vector Routing
- Each router advertises its distance vector every 30 seconds (or whenever its routing table changes) to all of its neighbors
- RIP always uses 1 as link metric
- Maximum hop count is 15, with “16” equal to “ $\infty$ ”
- Routes are timeout (set to 16) after 3 minutes if they are not updated

# RIP - History

- Late 1960s : Distance Vector protocols were used in the ARPANET
- Mid-1970s: XNS (Xerox Network system) routing protocol is the precursor of RIP in IP (and Novell's IPX RIP and Apple's routing protocol)
- 1982 Release of **routed** for BSD Unix
- 1988 RIPv1 (RFC 1058)
  - classful routing
- 1993 RIPv2 (RFC 1388)
  - adds subnet masks with each route entry
  - allows classless routing
- 1998 Current version of RIPv2 (RFC 2453)

# RIPv1 Packet Format



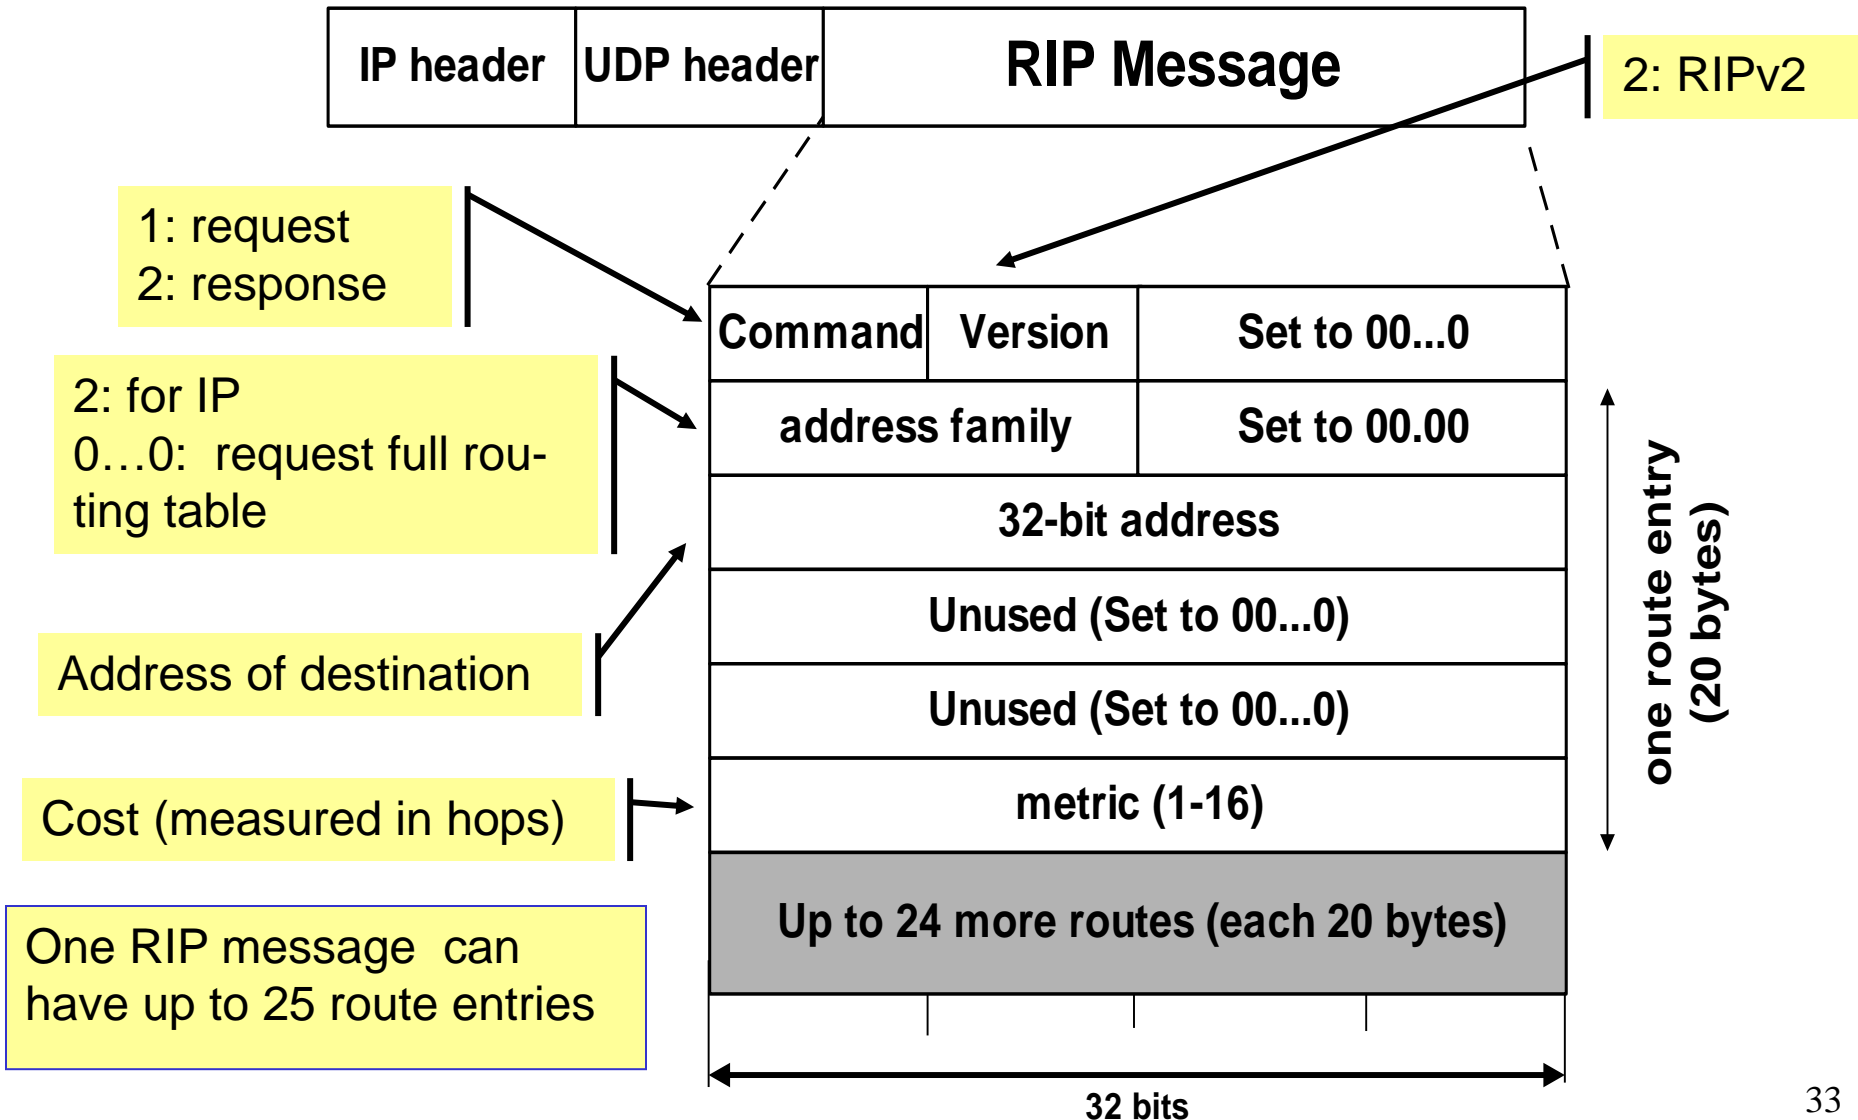
# RIPv2

---

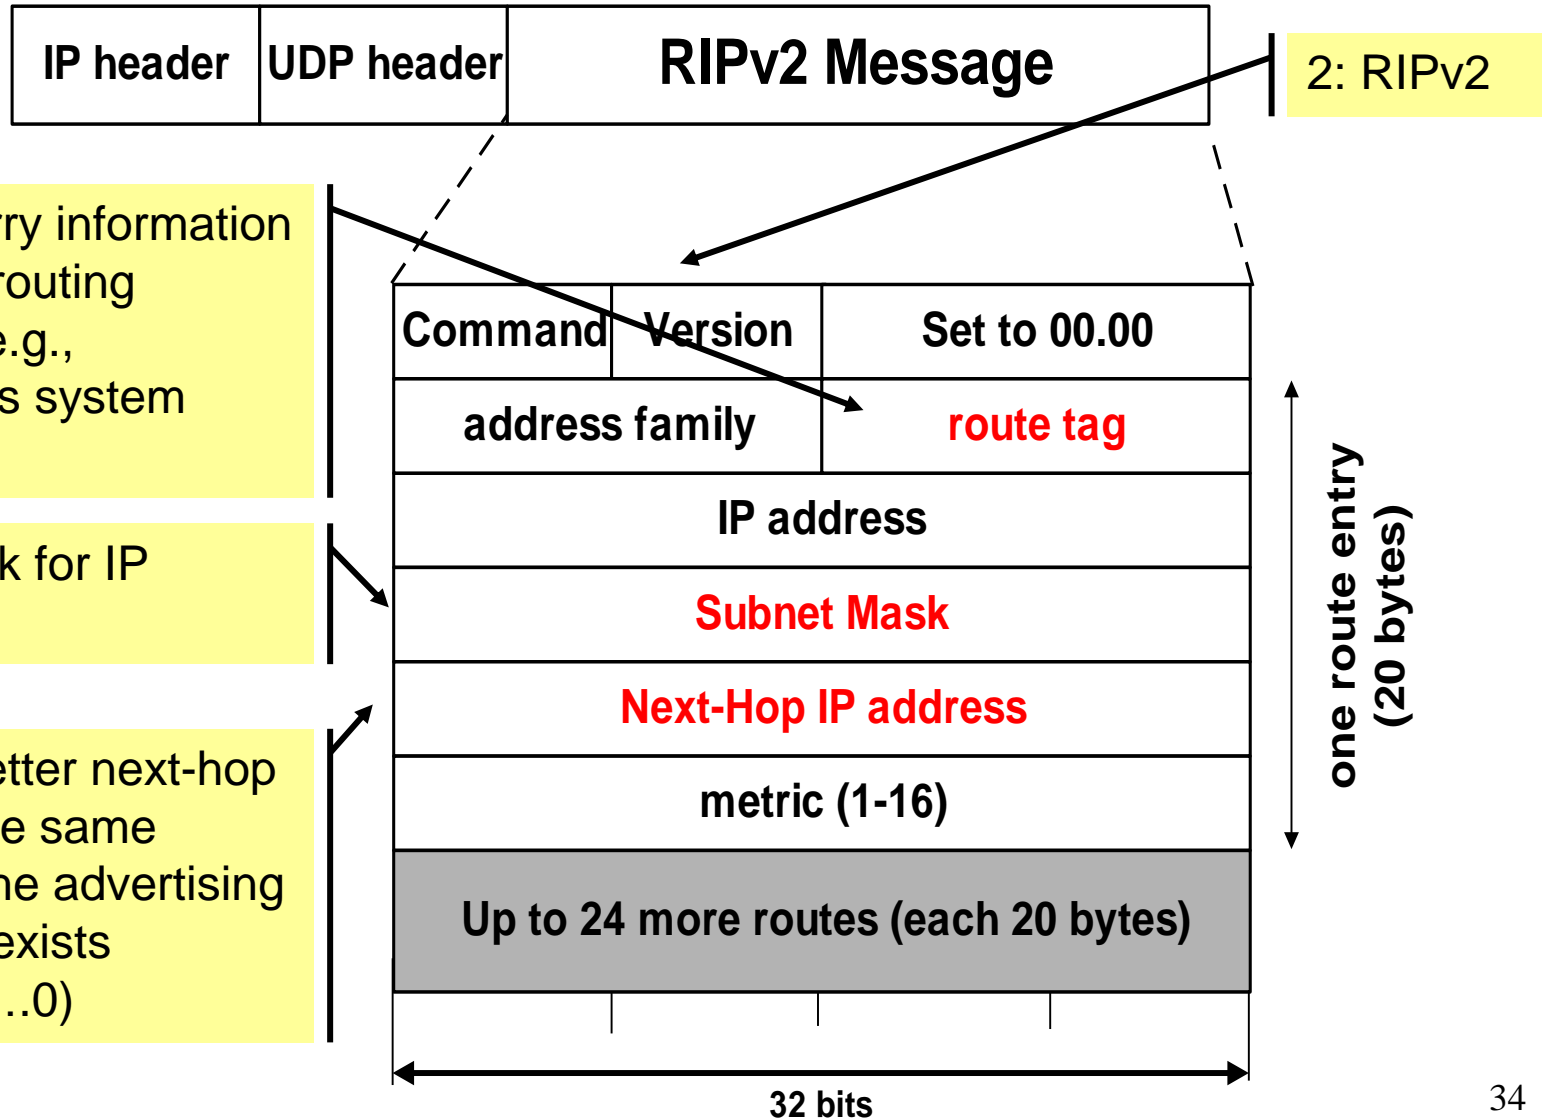
- RIPv2 is an extends RIPv1:
  - Subnet masks are carried in the route information
  - Authentication of routing messages
  - Route information carries next-hop address
  - Exploites IP multicasting
- Extensions of RIPv2 are carried in unused fields of RIPv1 messages



# RIPv2 Packet Format



# RIPv2 Packet Format



# RIP Messages

---

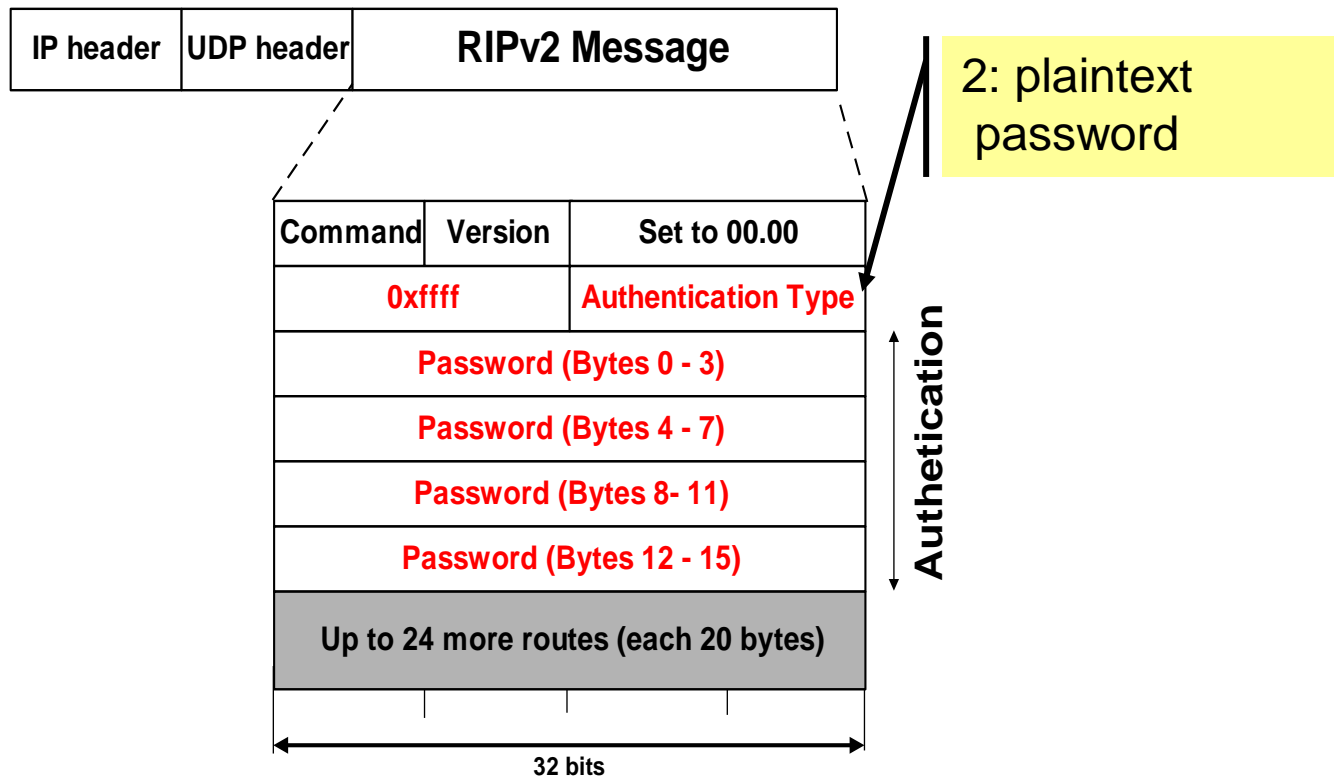
- This is the operation of RIP in **routed**. Dedicated port for RIP is UDP port 520.
- Two types of messages:
  - **Request messages**
    - used to ask neighboring nodes for an update
  - **Response messages**
    - contains an update

# Routing with RIP

- **Initialization:** Send a **request packet** (command = 1, address family=0..0) on all interfaces:
  - RIPv1 uses broadcast if possible,
  - RIPv2 uses multicast address 224.0.0.9, if possiblerequesting routing tables from neighboring routers
- **Request received:** Routers that receive above request send their entire routing table
- **Response received:** Update the routing table
- **Regular routing updates:** Every 30 seconds, send all or part of the routing tables to every neighbor in an response message
- **Triggered Updates:** Whenever the metric for a route change, send entire routing table.

# RIP Security

- Issue: Sending bogus routing updates to a router
- RIPv1: No protection
- RIPv2: Simple authentication scheme



# RIP Problems

---

- RIP takes a long time to stabilize
  - Even for a small network, it takes several minutes until the routing tables have settled after a change
- RIP has all the problems of distance vector algorithms, e.g., count-to-Infinity
  - » RIP uses split horizon to avoid count-to-infinity
- The maximum path in RIP is 15 hops