


TCP - Part I

Relates to Lab 5. First module on TCP which covers packet format, data transfer, and connection management.

Overview


TCP = Transmission Control Protocol

- Connection-oriented protocol
- Provides a reliable unicast end-to-end byte stream over an unreliable internetwork.



Connection-Oriented

- Before any data transfer, TCP establishes a **connection**:
 - One TCP entity is waiting for a connection (“**server**”)
 - The other TCP entity (“**client**”) contacts the server
- The actual procedure for setting up connections is more complex.
- Each connection is **CLIENT** full duplex




Reliable

- Byte stream is broken up into chunks which are called **segments**
 - Receiver sends acknowledgements (ACKs) for segments
 - TCP maintains a timer. If an ACK is not received in time, the segment is retransmitted
- **Detecting errors:**
 - TCP has checksums for header and data. Segments with invalid checksums are discarded
 - Each byte that is transmitted has a sequence number

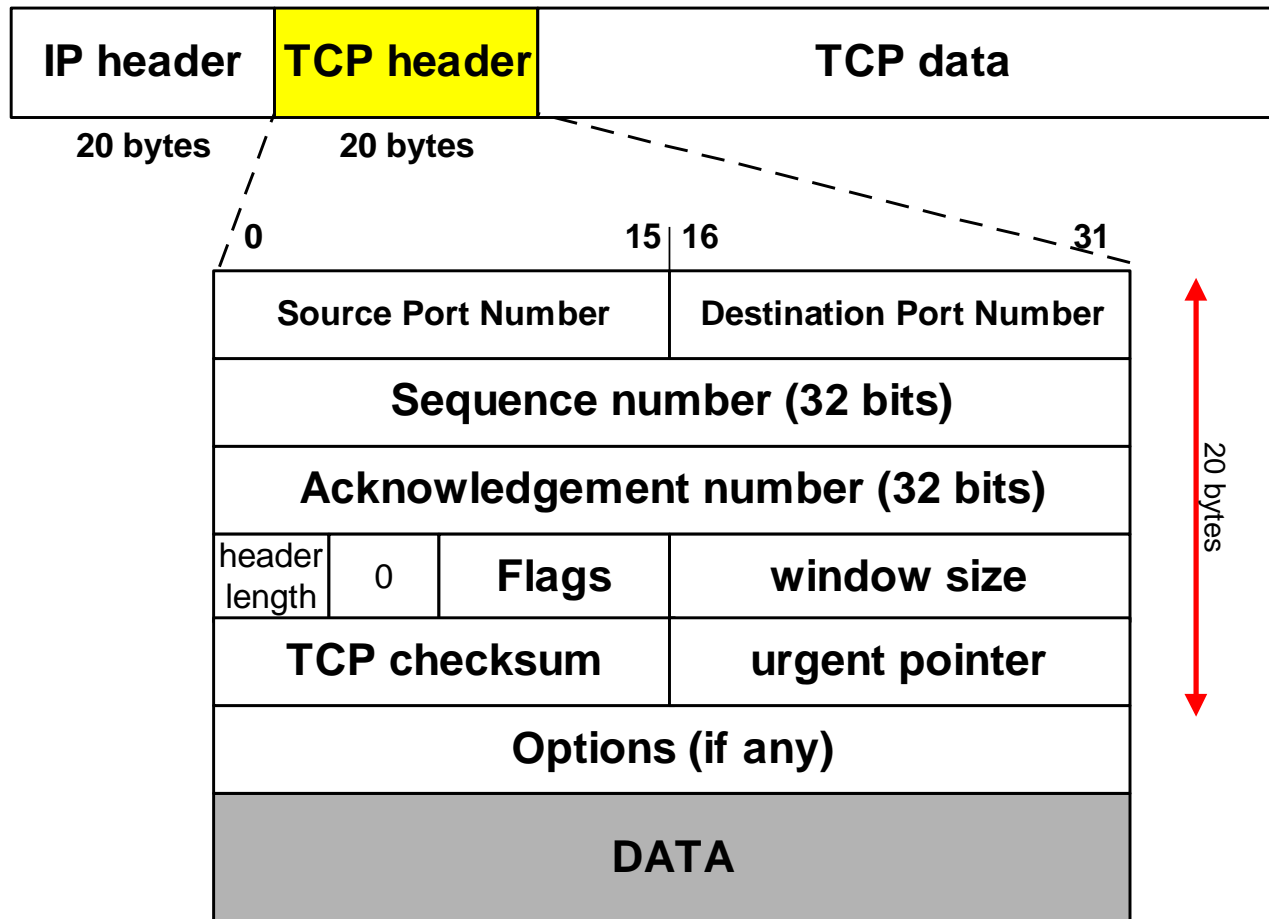
Byte Stream Service

- To the lower layers, TCP handles data in blocks, the segments.
- To the higher layers TCP handles data as a sequence of bytes and does not identify boundaries between bytes
- **So:** Higher layers do not know about the beginning and end of segments !



TCP Format


- TCP segments have a 20 byte header with ≥ 0 bytes of data.



TCP header fields

- **Port Number:**

- A port number identifies the endpoint of a connection.
- A pair **<IP address, port number>** identifies one endpoint of a connection.
- Two pairs **<client IP address, server port number>** and **<server IP address, server port number>** identify a TCP connection.



TCP header fields

- **Sequence Number (SeqNo):**

- Sequence number is 32 bits long.
- So the range of SeqNo is
$$0 \leq \text{SeqNo} \leq 2^{32} - 1 \approx 4.3 \text{ Gbyte}$$
- Each sequence number identifies a byte in the byte stream
- Initial Sequence Number (ISN) of a connection is set during connection establishment

Q: What are possible requirements for ISN ?

TCP header fields

- **Acknowledgement Number (AckNo):**

- Acknowledgements are piggybacked, I.e
a segment from A -> B can contain an
acknowledgement for a data sent in the B -> A direction

Q: Why is piggybacking good ?

- A hosts uses the AckNo field to send acknowledgements.
(If a host sends an AckNo in a segment it sets the “**ACK flag**”)
- The AckNo contains the next SeqNo that a hosts wants to
receive

Example: The acknowledgement for a segment with
sequence numbers 0-1500 is AckNo=1501

TCP header fields

- **Acknowledge Number (cont'd)**

- TCP uses the **sliding window flow protocol** (see CS 457) to regulate the flow of traffic from sender to receiver
- TCP uses the following variation of sliding window:
 - no NACKs (**N**egative **ACK**nowledgement)
 - only cumulative ACKs

- Example:

Assume: Sender sends two segments with “1..1500” and “1501..3000”, but receiver only gets the second segment.

In this case, the receiver cannot acknowledge the second packet. It can only send AckNo=1

TCP header fields

- **Header Length (4bits):**
 - Length of header in 32-bit words
 - Note that TCP header has variable length (with minimum 20 bytes)

TCP header fields

- **Flag bits:**

- **URG: Urgent pointer is valid**

- If the bit is set, the following bytes contain an urgent message in the range:

$\text{SeqNo} \leq \text{urgent message} \leq \text{SeqNo} + \text{urgent pointer}$

- **ACK: Acknowledgement Number is valid**

- **PSH: PUSH Flag**

- Notification from sender to the receiver that the receiver should pass all data that it has to the application.
- Normally set by sender when the sender's buffer is empty

TCP header fields

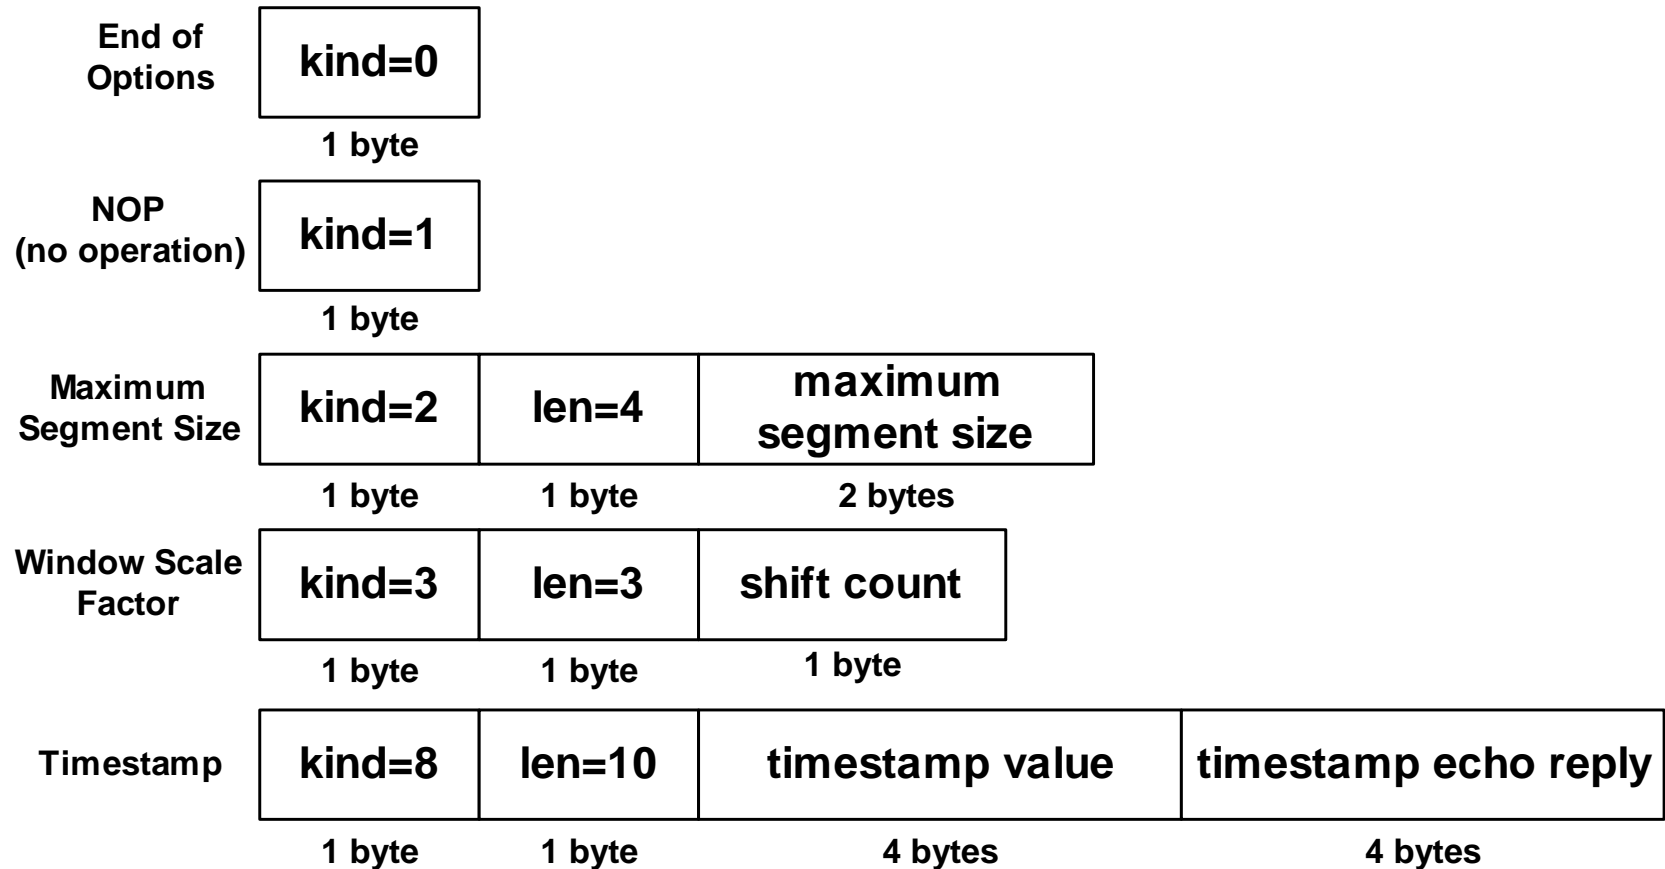
- **Flag bits:**
 - **RST: Reset the connection**
 - The flag causes the receiver to reset the connection
 - Receiver of a RST terminates the connection and indicates higher layer application about the reset
 - **SYN: Synchronize sequence numbers**
 - Sent in the first packet when initiating a connection
 - **FIN: Sender is finished with sending**
 - Used for closing a connection
 - Both sides of a connection must send a **FIN**

TCP header fields

- **Window Size:**
 - Each side of the connection advertises the window size
 - Window size is the maximum number of bytes that a receiver can accept.
 - Maximum window size is $2^{16}-1=65535$ bytes
- **TCP Checksum:**
 - TCP checksum covers over both TCP header **and** TCP data (also covers some parts of the IP header)
- **Urgent Pointer:**
 - Only valid if **URG** flag is set

TCP header fields

- Options:



TCP header fields

- **Options:**
 - **NOP** is used to pad TCP header to multiples of 4 bytes
 - **Maximum Segment Size**
 - **Window Scale Options**
 - » Increases the TCP window from 16 to 32 bits, I.e., the window size is interpreted differently
 - Q: What is the different interpretation ?*
 - » This option can only be used in the SYN segment (first segment) during connection establishment time
 - **Timestamp Option**
 - » Can be used for roundtrip measurements

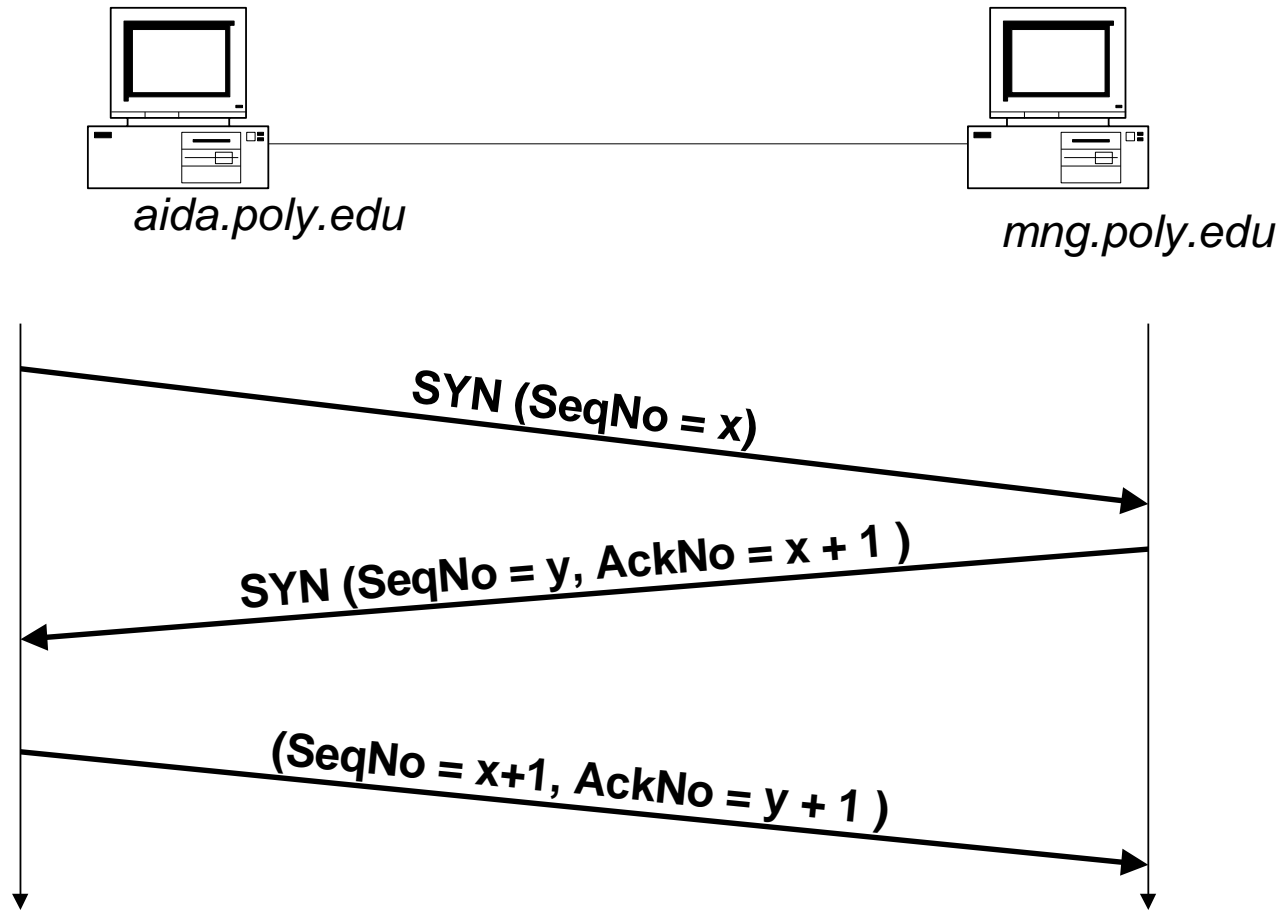
Connection Management in TCP

- **Opening a TCP Connection**
- **Closing a TCP Connection**
- **Special Scenarios**
- **State Diagram**

TCP Connection Establishment


- TCP uses a **three-way handshake** to open a connection:
 - (1) ACTIVE OPEN:** Client sends a segment with
 - SYN bit set *
 - port number of client
 - initial sequence number (ISN) of client
 - (2) PASSIVE OPEN:** Server responds with a segment with
 - SYN bit set *
 - initial sequence number of server
 - ACK for ISN of client
 - (3) Client acknowledges by sending a segment with:**
 - ACK ISN of server
- (* counts as one byte)

Three-Way Handshake




A Closer Look with tcpdump

aida issues
an "telnet mng"



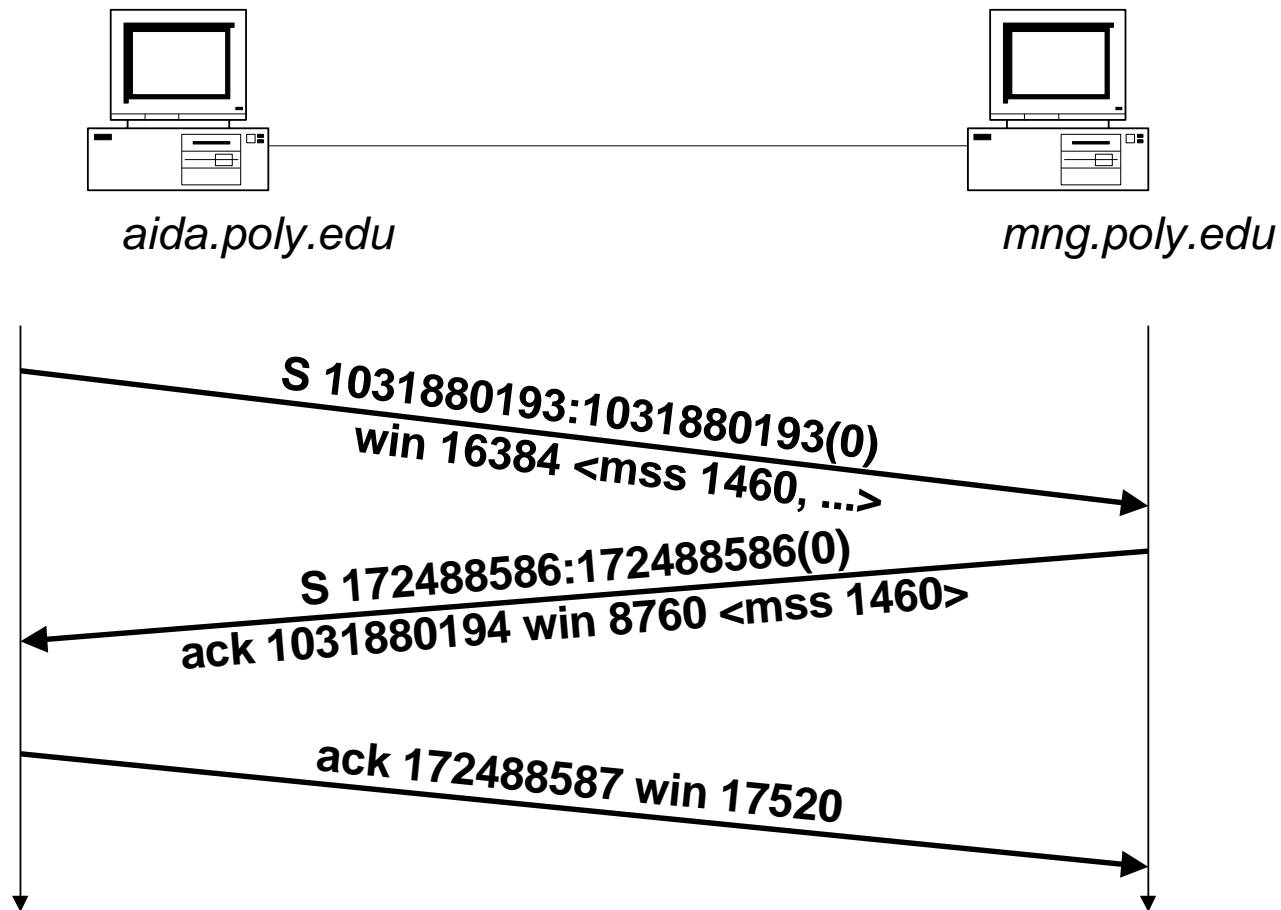
aida.poly.edu



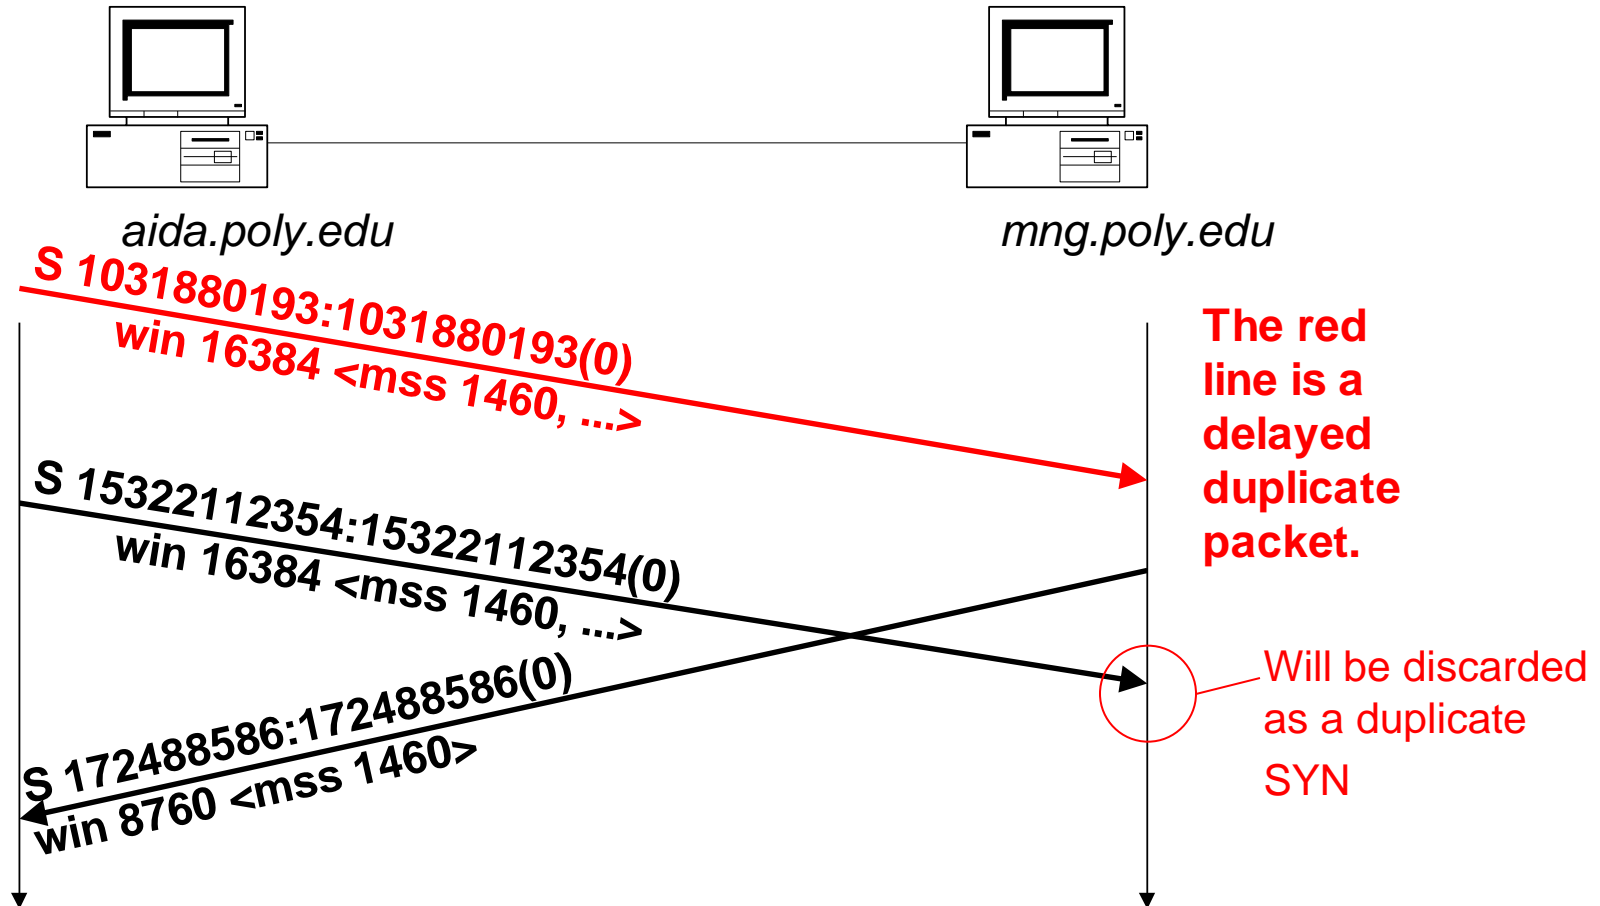
mng.poly.edu

- 1 aida.poly.edu.1121 > mng.poly.edu.telnet: S 1031880193:1031880193(0)
win 16384 <mss 1460,nop,wscale 0,nop,nop,timestamp>
- 2 mng.poly.edu.telnet > aida.poly.edu.1121: S 172488586:172488586(0)
ack 1031880194 win 8760 <mss 1460>
- 3 aida.poly.edu.1121 > mng.poly.edu.telnet: . ack 172488587 win 17520
- 4 aida.poly.edu.1121 > mng.poly.edu.telnet: P 1031880194:1031880218(24)
ack 172488587 win 17520
- 5 mng.poly.edu.telnet > aida.poly.edu.1121: P 172488587:172488590(3)
ack 1031880218 win 8736
- 6 aida.poly.edu.1121 > mng.poly.edu.telnet: P 1031880218:1031880221(3)
ack 172488590 win 17520

Three-Way Handshake



Why is a Two-Way Handshake not enough?




When aida initiates the data transfer (starting with SeqNo=15322112355), mng will reject all data.

TCP Connection Termination

- Each end of the data flow must be shut down independently (**“half-close”**)
- If one end is done it sends a FIN segment. This means that no more data will be sent
- Four steps involved:
 - (1) X sends a FIN to Y (**active close**)
 - (2) Y ACKs the FIN,
(at this time: Y can still send data to X)
 - (3) and Y sends a FIN to X (**passive close**)
 - (4) X ACKs the FIN.

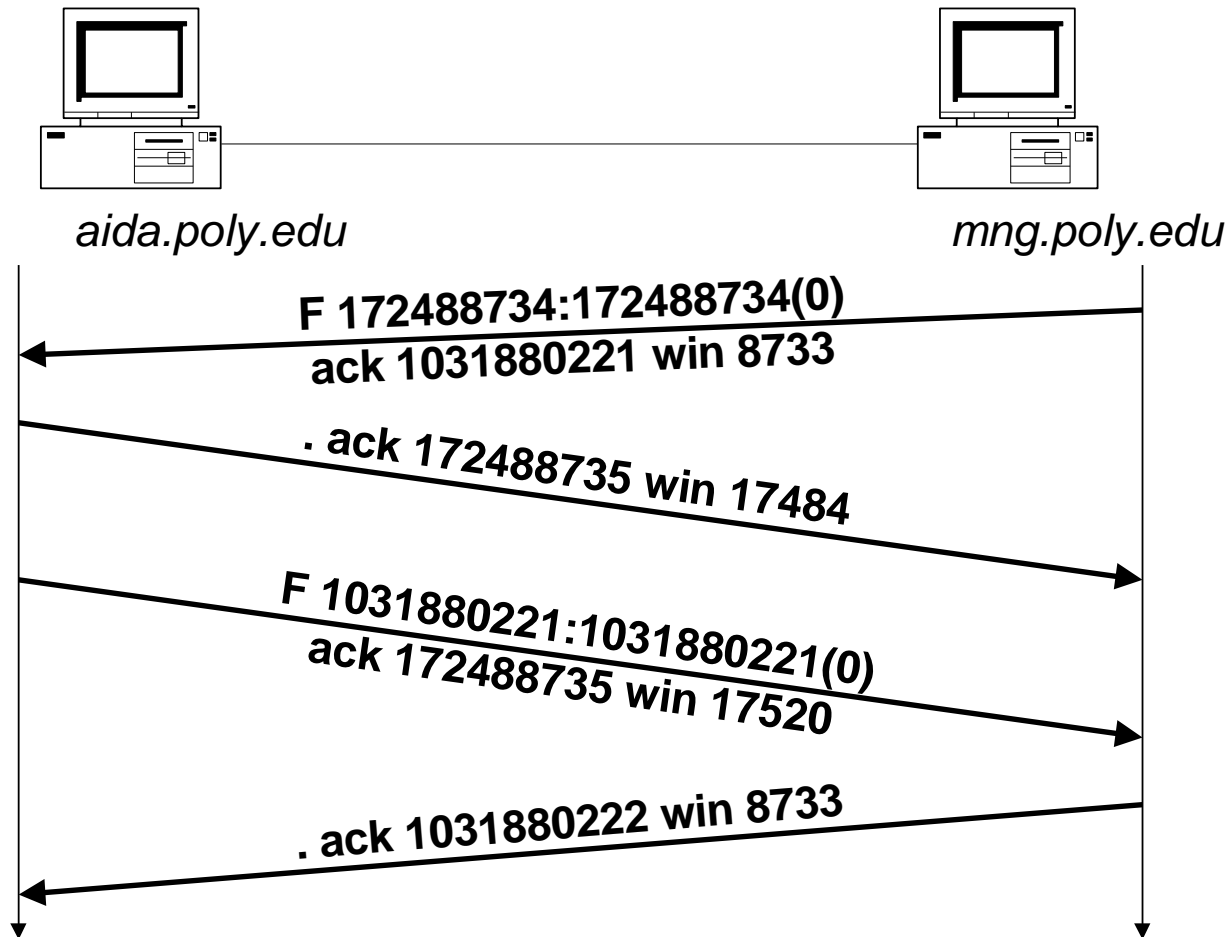
Connection termination with tcpdump

aida issues
an "telnet mng"



- 1 mng.poly.edu.telnet > aida.poly.edu.1121: F 172488734:172488734(0)
ack 1031880221 win 8733
- 2 aida.poly.edu.1121 > mng.poly.edu.telnet: . ack 172488735 win 17484
- 3 aida.poly.edu.1121 > mng.poly.edu.telnet: F 1031880221:1031880221(0)
ack 172488735 win 17520
- 4 mng.poly.edu.telnet > aida.poly.edu.1121: . ack 1031880222 win 8733

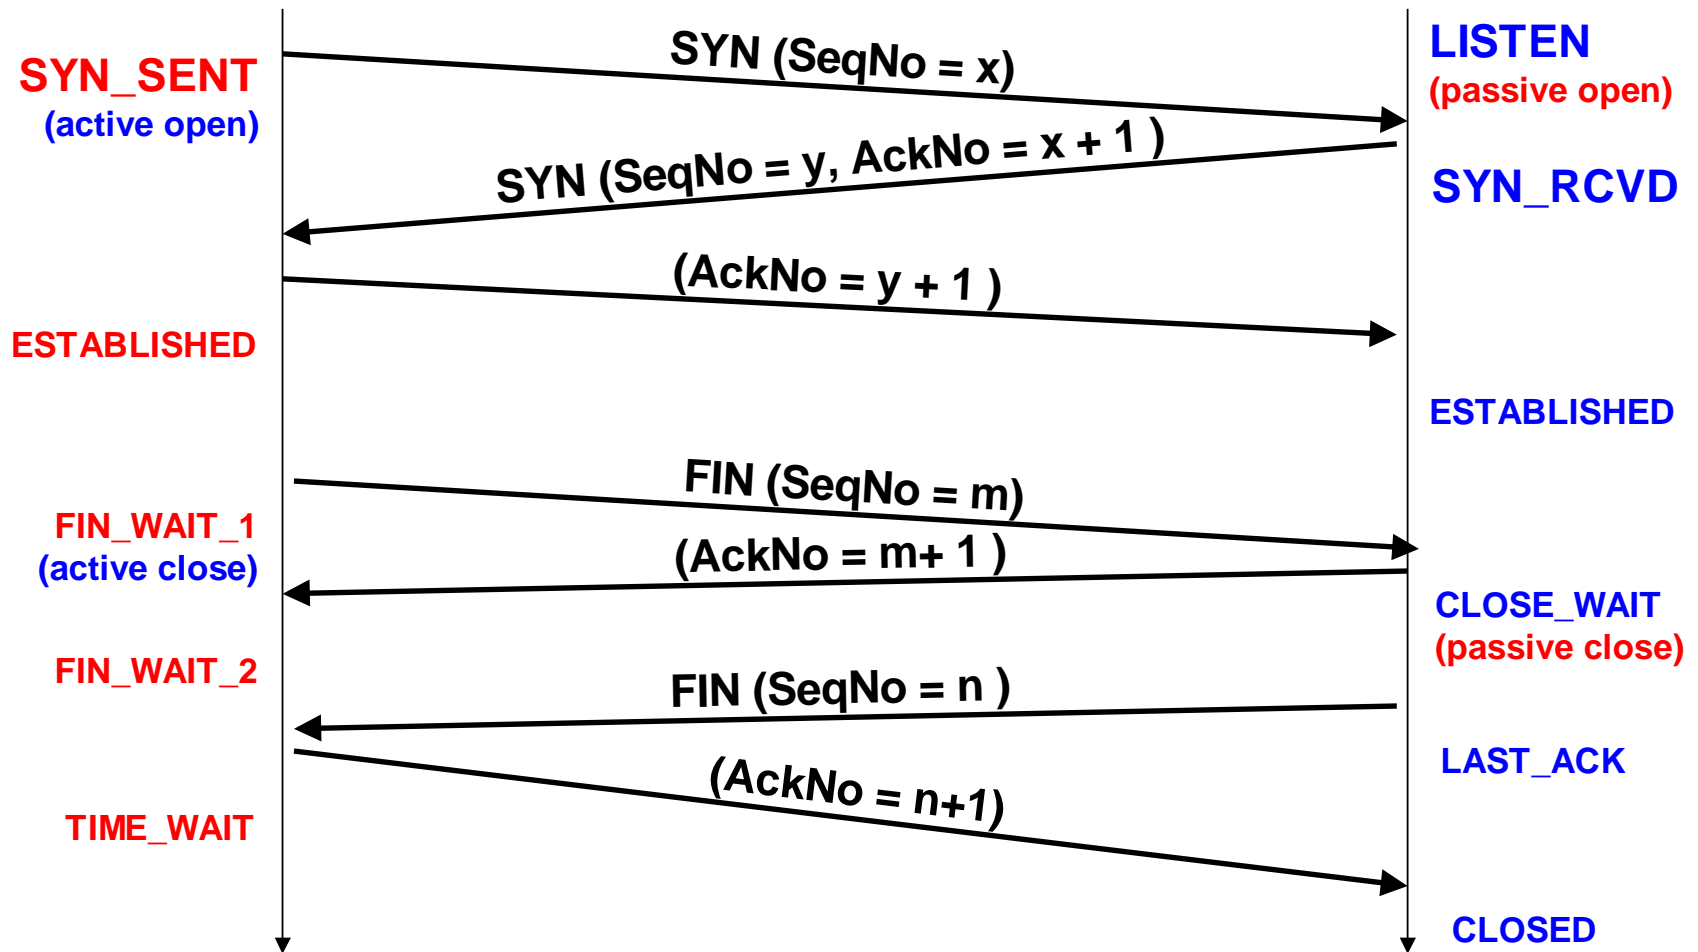
TCP Connection Termination



TCP States

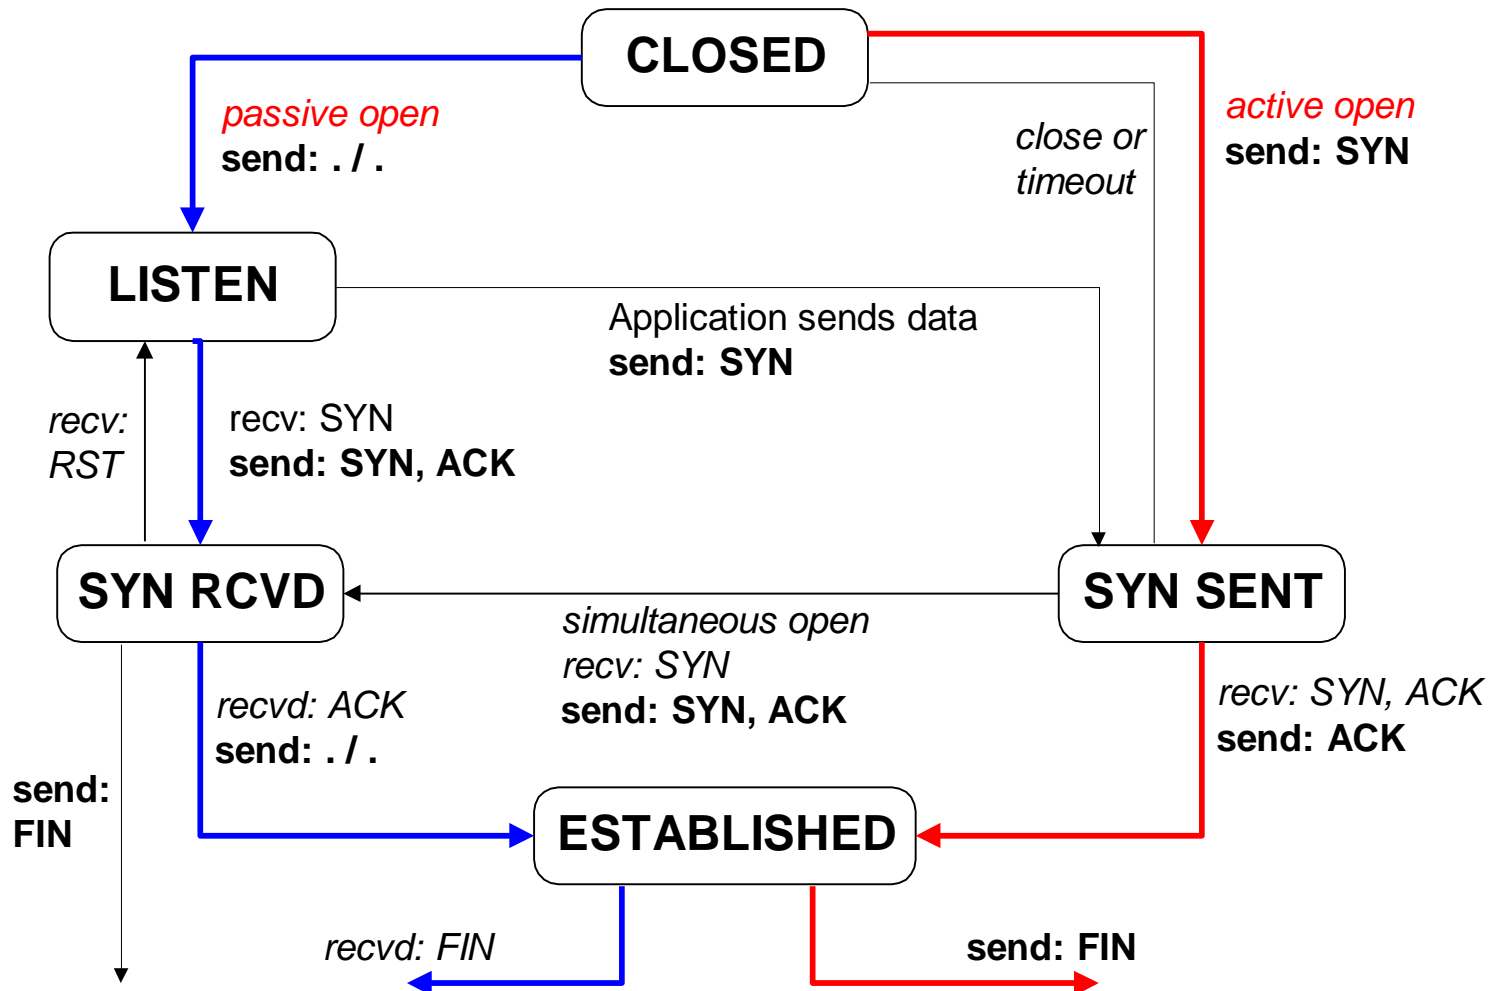
State	Description
CLOSED	No connection is active or pending
LISTEN	The server is waiting for an incoming call
SYN RCVD	A connection request has arrived; wait for Ack
SYN SENT	The client has started to open a connection
ESTABLISHED	Normal data transfer state
FIN WAIT 1	Client has said it is finished
FIN WAIT 2	Server has agreed to release
TIMED WAIT	Wait for pending packets (“2MSL wait state”)
CLOSING	Both Sides have tried to close simultanesously
CLOSE WAIT	Server has initiated a release
LAST ACK	Wait for pending packets

TCP States in “Normal” Connection Lifetime



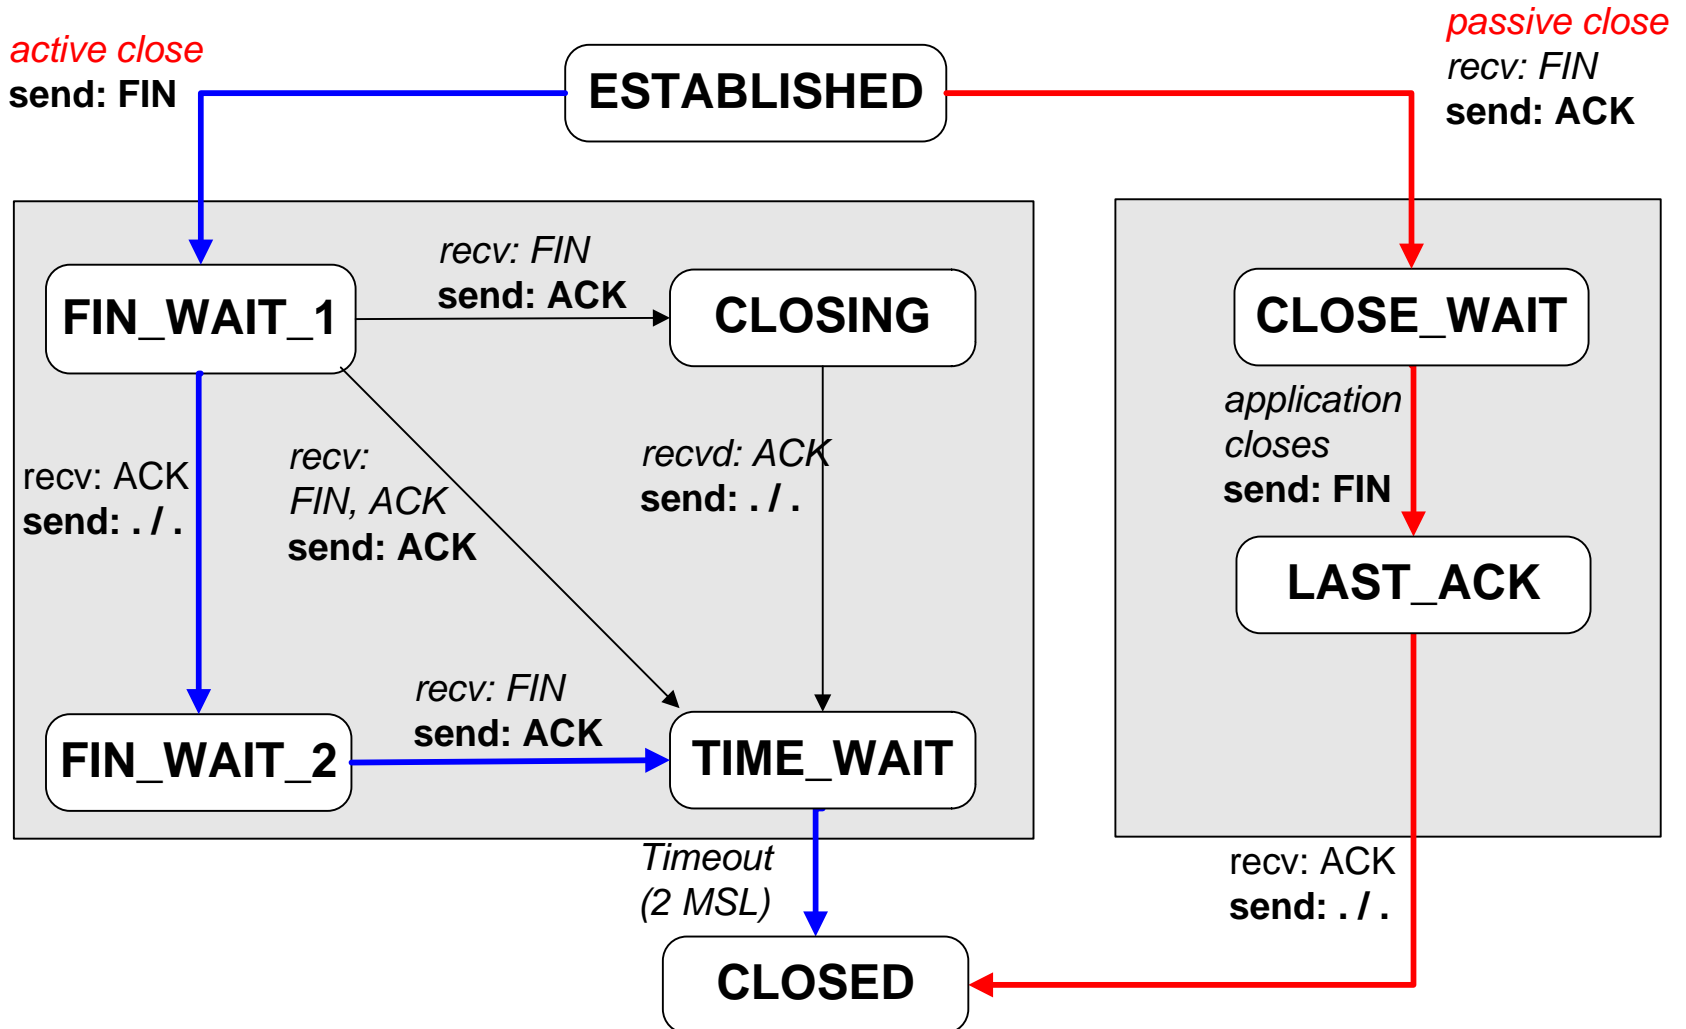
TCP State Transition Diagram

Opening A Connection



TCP State Transition Diagram

Closing A Connection



2MSL Wait State

2MSL Wait State = TIME_WAIT

- When TCP does an active close, and sends the final ACK, the connection **must stay in in the TIME_WAIT state for twice the maximum segment lifetime.**

2MSL= 2 * Maximum Segment Lifetime

- Why?
TCP is given a chance to resent the final ACK. (Server will timeout after sending the FIN segment and resend the FIN)
- The MSL is set to 2 minutes or 1 minute or 30 seconds.

Resetting Connections

- Resetting connections is done by setting the RST flag
- **When is the RST flag set?**
 - Connection request arrives and no server process is waiting on the destination port
 - Abort (Terminate) a connection
Causes the receiver to throw away buffered data. Receiver does not acknowledge the RST segment