

# TCP III - Error Control

---

TCP Error Control

# ARQ Error Control

---

- Two types of errors:
  - **Lost packets**
  - **Damaged packets**
- Most Error Control techniques are based on:
  1. Error Detection Scheme (Parity checks, CRC).
  2. Retransmission Scheme.
- Error control schemes that involve error detection and retransmission of lost or corrupted packets are referred to as **Automatic Repeat Request (ARQ) error control**.


# Background: ARQ Error Control

---

- The most common ARQ retransmission schemes:
  - Stop-and-Wait ARQ
  - Go-Back-N ARQ
  - Selective Repeat ARQ
- The protocol for sending ACKs in all ARQ protocols are based on the sliding window flow control scheme
- TCP uses a version of the Go-Back-N Protocol

# Background: Go-Back-N ARQ

- Go-Back-N sends cumulative acknowledgments



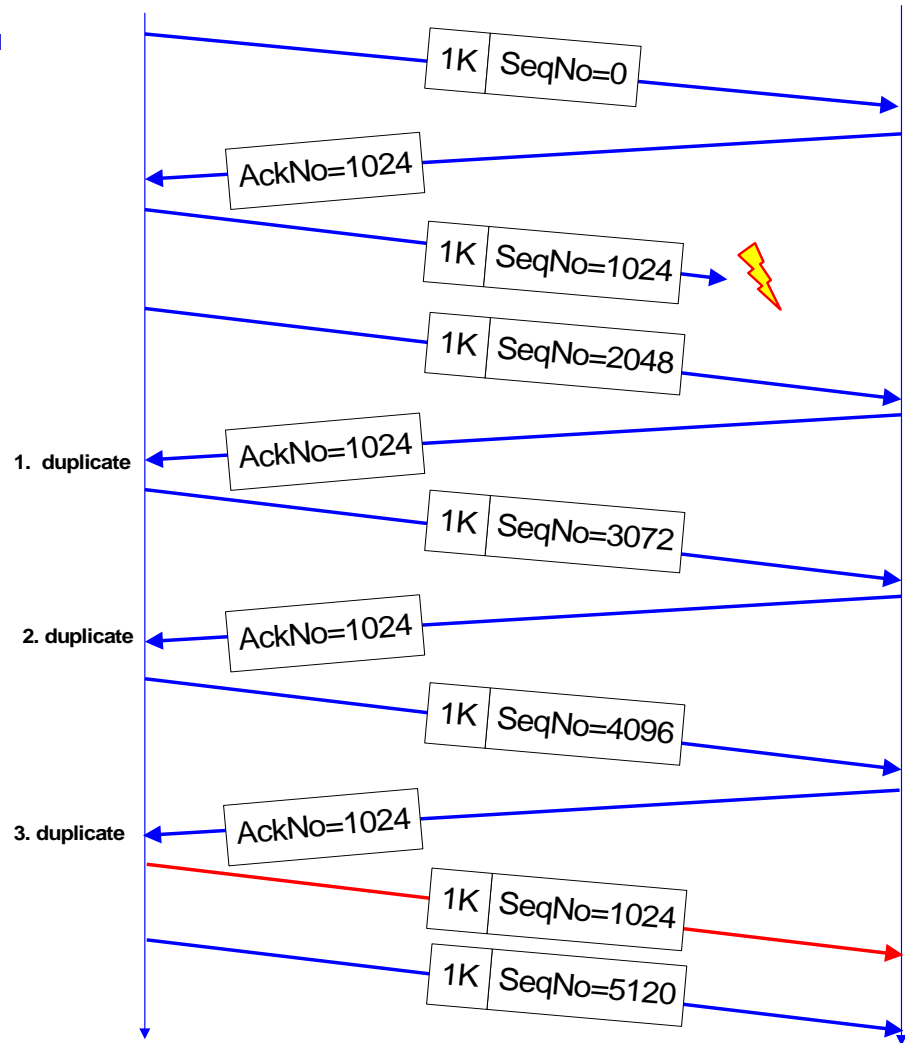
# Retransmissions in TCP

---

- A TCP sender retransmits a segment when it assumes that the segment has been lost:
  1. No ACK has been received and a timeout occurs
  2. Multiple ACKs have been received for the same segment

# Receiving duplicate ACKs

- If three or more duplicate ACKs are received in a row, the TCP sender believes that a segment has been lost.
- Then TCP performs a retransmission of what seems to be the missing segment, without waiting for a timeout to happen.
- This can fix losses of single segments



# Retransmission Timer

- TCP sender maintains one retransmission timer for each connection
- When the timer reaches the retransmission timeout (RTO) value, the sender retransmits the first segment that has not been acknowledged
- The timer is started when
  1. When a packet with payload is transmitted and timer is not running
  2. When an ACK arrives that acknowledges new data,
  3. When a segment is retransmitted
- The timer is stopped when
  - All segments are acknowledged

# How to set the timer

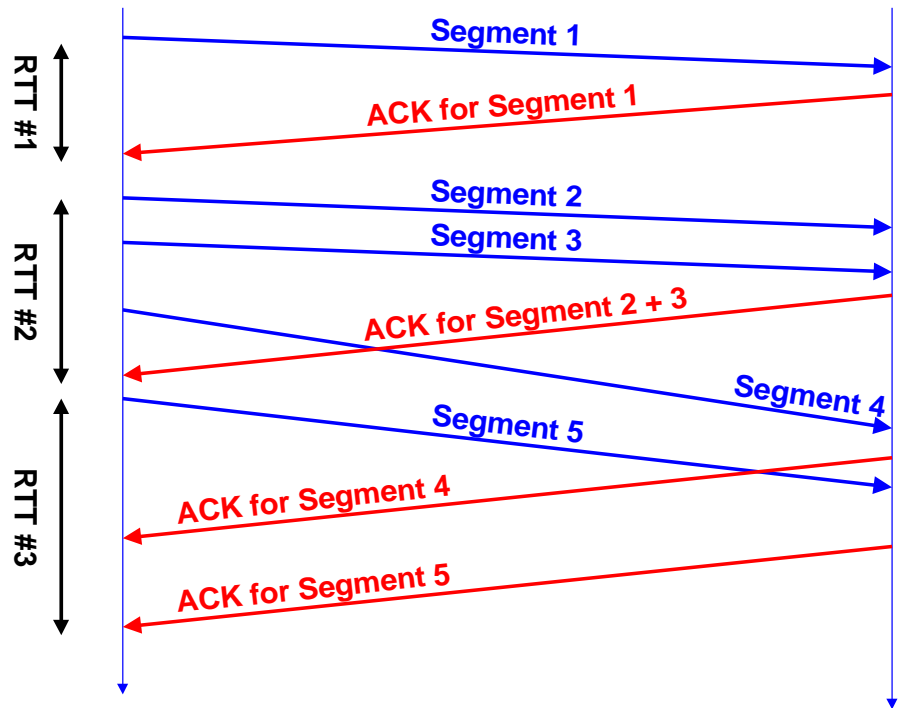
- **Retransmission Timer:**

- The setting of the retransmission timer is crucial for good performance of TCP
- **Timeout value too small** → results in unnecessary retransmissions
- **Timeout value too large** → long waiting time before a retransmission can be issued
- A problem is that the delays in the network are not fixed
- Therefore, the retransmission timers must be adaptive



# Setting the value of RTO:

- The RTO value is set based on round-trip time (RTT) measurements that each TCP performs
- Each TCP connection measures the time difference between the transmission of a segment and the receipt of the corresponding ACK
- There is only one measurement ongoing at any time (i.e., measurements do not overlap)
- Figure on the right shows three RTT measurements



# Setting the RTO value

- RTO is calculated based on the *RTT* measurements
  - Uses an exponential moving average to calculate estimators for delay (*srtt*) and variance of delay (*rttvar*) from
- The RTT measurements are smoothed by the following estimators *srtt* and *rttvar*:

$$\begin{aligned} srtt_{n+1} &= \alpha \text{ RTT} + (1 - \alpha) srtt_n \\ rttvar_{n+1} &= \beta ( | \text{RTT} - srtt_n | ) + (1 - \beta) rttvar_n \\ RTO_{n+1} &= srtt_{n+1} + 4 rttvar_{n+1} \end{aligned}$$

- The gains are set to  $\alpha = 1/4$  and  $\beta = 1/8$

# Setting the RTO value (cont'd)

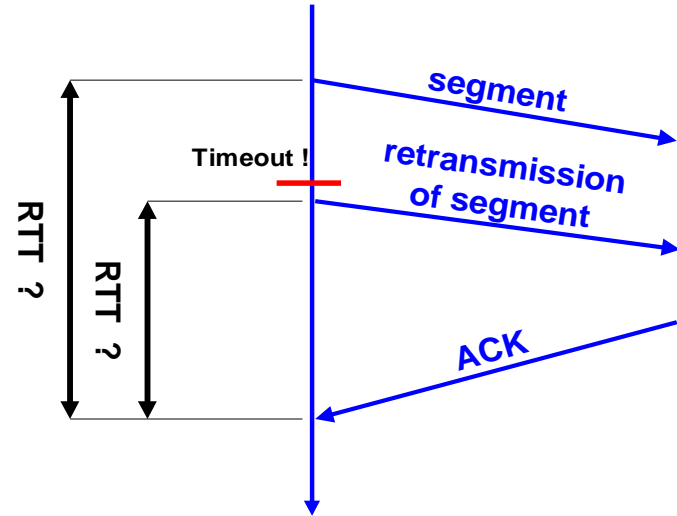
- Initial value for RTO:
  - Sender should set the initial value of RTO to
$$RTO_0 = 3 \text{ seconds}$$
- RTO calculation after first RTT measurements arrived
$$srtt_1 = RTT$$
$$rttvar_1 = RTT / 2$$
$$RTO_1 = srtt_1 + 4 rttvar_{n+1}$$
- When a timeout occurs, the RTO value is doubled
$$RTO_{n+1} = \max ( 2 RTO_n, 64 ) \text{ seconds}$$

This is called *an exponential backoff*

# Karn's Algorithm

If an ACK for a retransmitted segment is received, the sender cannot tell if the ACK belongs to the original or the retransmission.

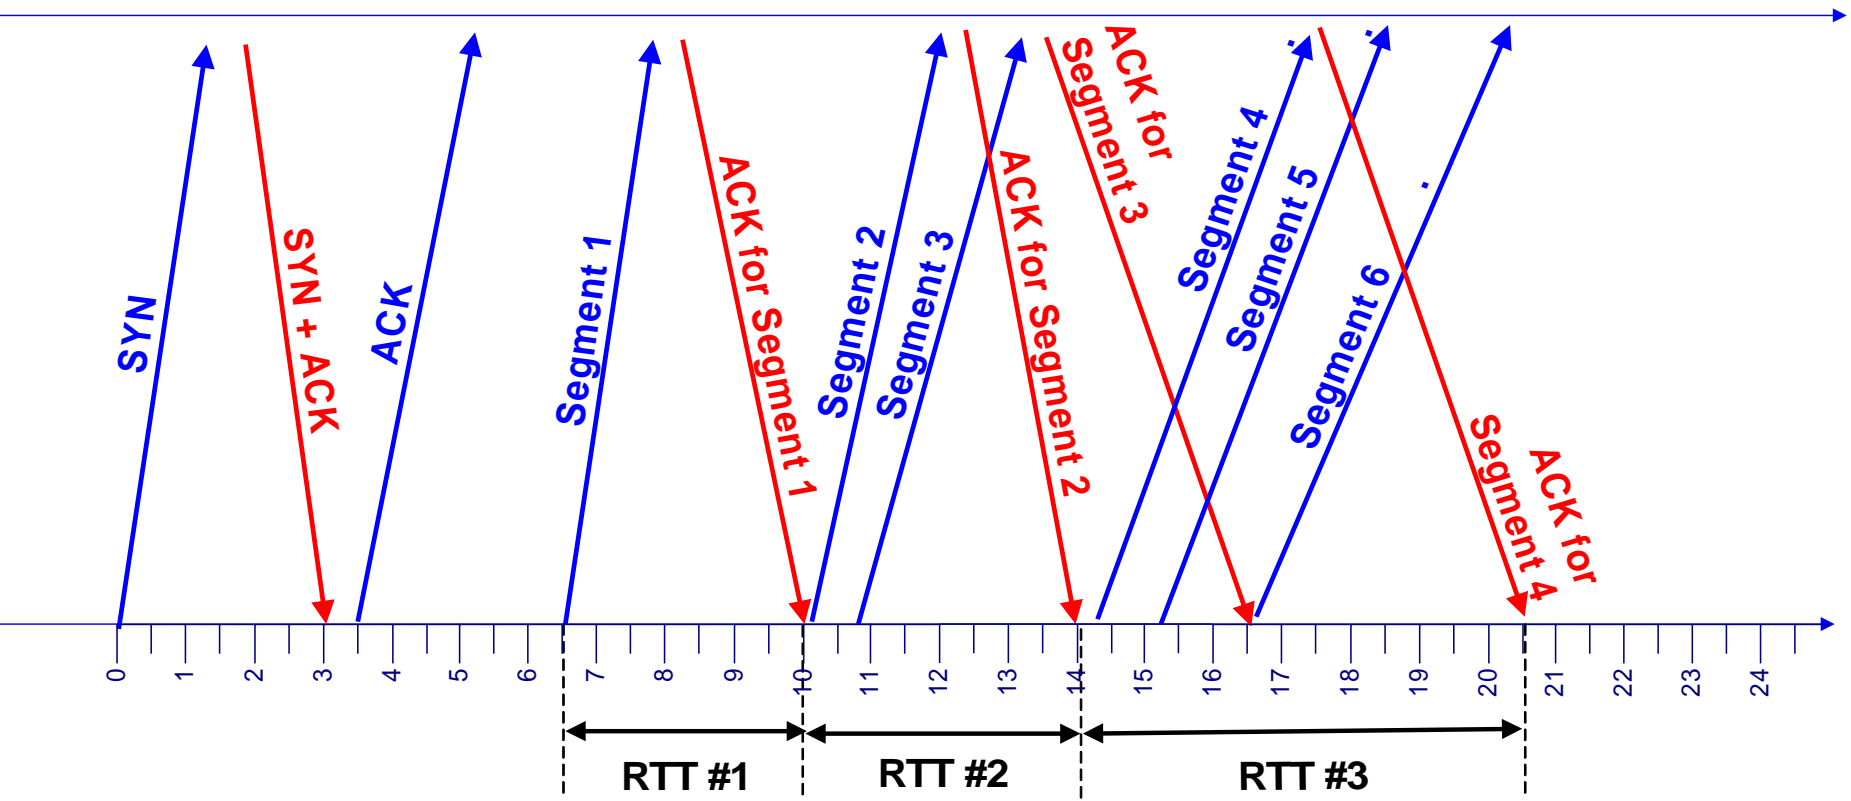
→ RTT measurements is ambiguous in this case



## Karn's Algorithm:

- Don't update *RTT* on any segments that have been retransmitted
- Restart RTT measurements only after an ACK is received for a segment that is not retransmitted

# RTO Calculation: Example



# TCP Retransmission Timer

- **Retransmission Timer:**

- The setting of the retransmission timer is crucial for efficiency
- **Timeout value too small** → results in unnecessary retransmissions
- **Timeout value too large** → long waiting time before a retransmission can be issued
- A problem is that the delays in the network are not fixed
- Therefore, the retransmission timers must be adaptive

# Round-Trip Time Measurements

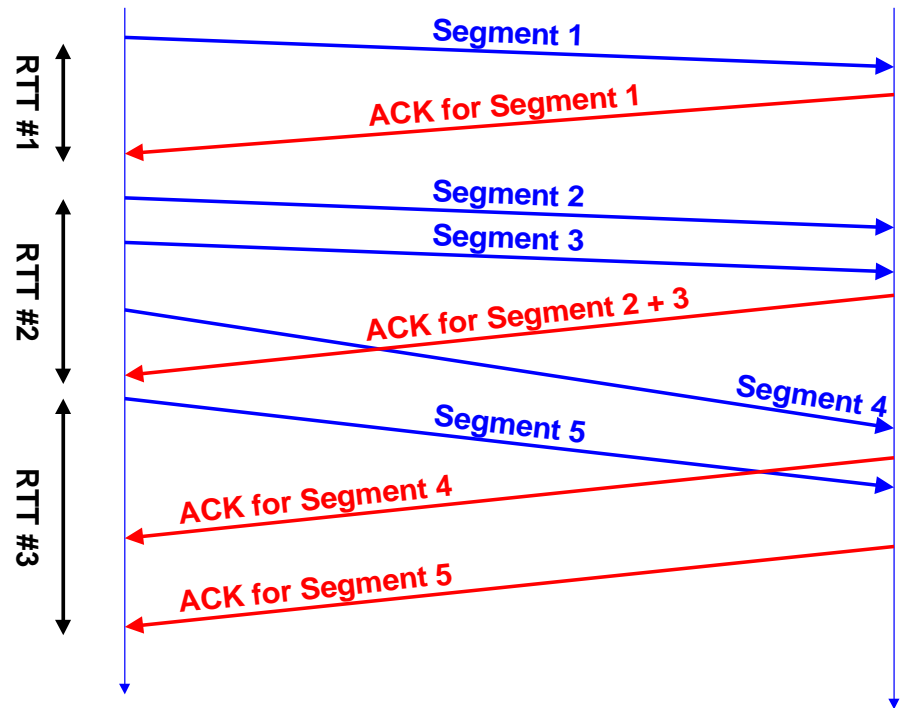
- The retransmission mechanism of TCP is adaptive
- The retransmission timers are set based on round-trip time (RTT) measurements that TCP performs

The RTT is based on time difference between segment transmission and ACK

**But:**

TCP does not ACK each segment

Each connection has only one timer



# Round-Trip Time Measurements

- Retransmission timer is set to a **Retransmission Timeout (RTO) value**.
- RTO is calculated based on the *RTT* measurements.
- The RTT measurements are smoothed by the following estimators *srtt* and *rttvar*:

$$srtt_{n+1} = \alpha \text{ RTT} + (1 - \alpha) srtt_n$$

$$rttvar_{n+1} = \beta ( | \text{RTT} - srtt_{n+1} | ) + (1 - \beta) rttvar_n$$


$$RTO_{n+1} = srtt_{n+1} + 4 rttvar_{n+1}$$

- The gains are set to  $\alpha = 1/4$  and  $\beta = 1/8$
- $srtt_0 = 0 \text{ sec}$ ,  $rttvar_0 = 3 \text{ sec}$ , Also:  $RTO_1 = srtt_1 + 2 rttvar_1$



# Karn's Algorithm

- If an ACK for a retransmitted segment is received, the sender cannot tell if the ACK belongs to the original or the retransmission.




## Karn's Algorithm:

Don't update *srtt* on any segments that have been retransmitted.  
Each time when TCP retransmits, it sets:

$$RTO_{n+1} = \max(2 RTO_n, 64) \quad (\text{exponential backoff})$$

# Measuring TCP Retransmission Timers


*argon.tcpip-lab.edu*  
("Argon")



Web client

Transfer file

*neon.tcpip-lab.edu*  
("Neon")



Web server

- Transfer file from Argon to neon
- Unplug Ethernet of Argon cable in the middle of file transfer

# Interpreting the Measurements

- The interval between retransmission attempts in seconds is:  
*1.03, 3, 6, 12, 24, 48, 64, 64, 64, 64, 64, 64.*
- Time between retransmissions is doubled each time (**Exponential Backoff Algorithm**)
- Timer is not increased beyond 64 seconds
- TCP gives up after 13th attempt and 9 minutes.

