

IP Multicasting

Relates to Lab 10.

It covers IP multicasting, including multicast addressing, IGMP, and multicast routing.

Applications with multiple receivers

- Many applications transmit the same data at one time to multiple receivers
 - Broadcasts of Radio or Video
 - Videoconferencing
 - Shared Applications
- A network must have mechanisms to support such applications in an efficient manner


Motivation

"Together, Internet broadcasting and multicasting are the next chapters in the evolution of the Internet as a revolutionary catalyst for the information age."

*Vint Cerf, Senior vice president of
MCI/Worldcom, April 1999.*

Multicasting


- Multicast communications refers to one-to-many or many-to-many communications.



IP Multicasting refers to the implementation of multicast communication in the Internet


Multicasting over a Packet Network

- Without support for multicast at the network layer:



Multicasting over a Packet Network

- With support for multicast at the network layer:



- **Requires a set of mechanisms:**
 - (1) Packet forwarding can send multiple copies of same packet
 - (2) Multicast routing algorithm which builds a spanning tree (dynamically)

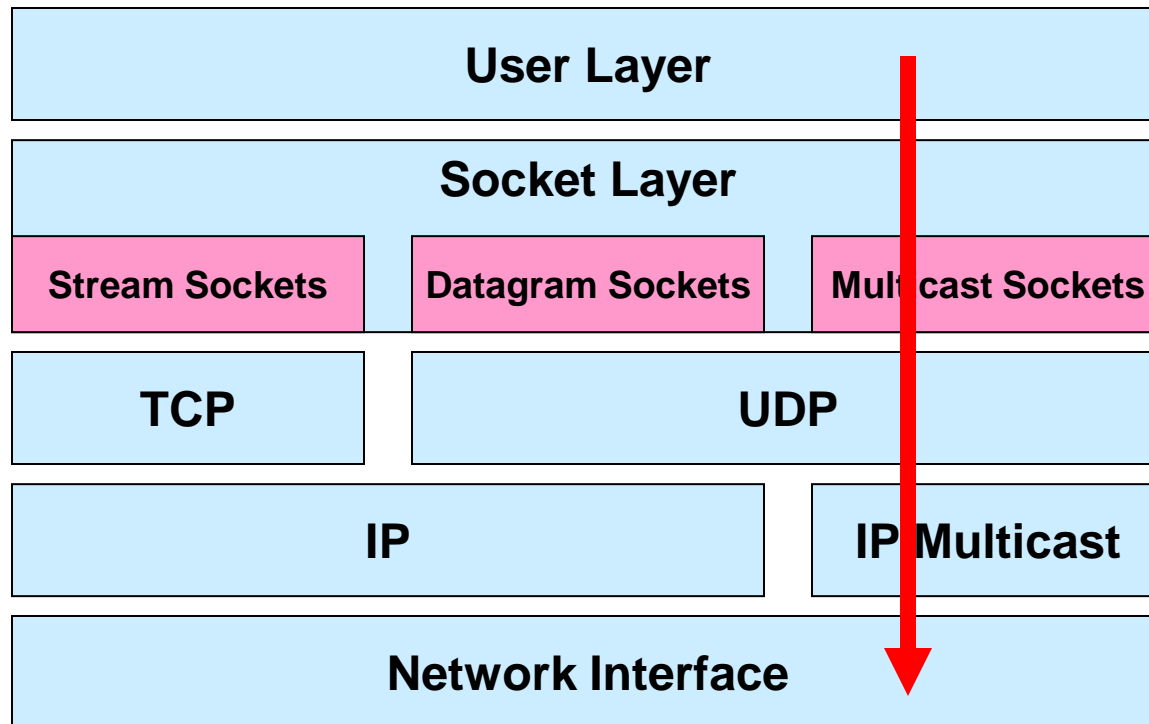
Semantics of IP Multicast

IP multicast works as follows:

- Multicast groups are identified by IP addresses in the range 224.0.0.0 - 239.255.255.255 (class D address)
- Every host (*more precisely*: interface) can join and leave a multicast group dynamically
 - » no access control
- Every IP datagram send to a multicast group is transmitted to all members of the group
 - » no security, no “floor control”
- The IP Multicast service is unreliable

The IP Protocol Stack

- IP Multicasting only supports UDP as higher layer
- There is no multicast TCP !



IP Multicasting

- There are three essential components of the IP Multicast service:

IP Multicast Addressing
IP Group Management
Multicast Routing

Multicast Addressing

- All Class D addresses are multicast addresses:

Class D

1	1	1	0	multicast group id
28 bits				

Class	From	To
D	224.0.0.0	239.255.255.255


- Multicast addresses are dynamically assigned.
- An IP datagram sent to a multicast address is forwarded to everyone who has joined the multicast group
- If an application is terminated, the multicast address is (implicitly) released.

Types of Multicast addresses

- The range of addresses between 224.0.0.0 and 224.0.0.255, inclusive, is reserved for the use of routing protocols and other low-level topology discovery or maintenance protocols
- Multicast routers should not forward any multicast datagram with destination addresses in this range.
- Examples of special and reserved Class D addresses, e.g,
 - 224.0.0.1** All systems on this subnet
 - 224.0.0.2** All routers on this subnet
 - 224.0.1.1** NTP (Network Time Protocol)
 - 224.0.0.9** RIP-2 (a routing protocol)

Multicast Address Translation


- In Ethernet MAC addresses, a multicast address is identified by setting the lowest bit of the “most left byte”



Not all Ethernet cards can filter multicast addresses in hardware

- Then: Filtering is done in software by device driver.

Multicast Address Mapping



IGMP

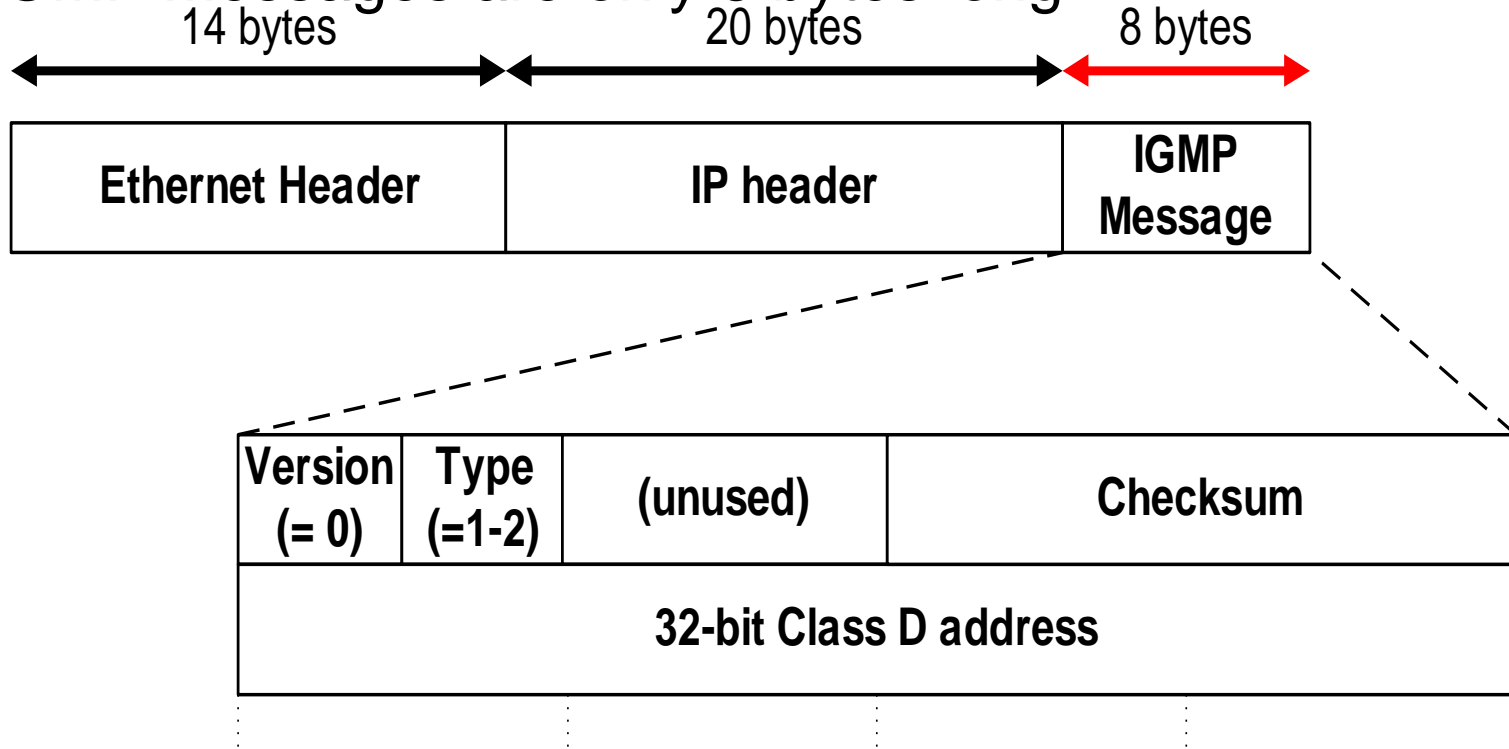
- The **Internet Group Management Protocol (IGMP)** is a simple protocol for the support of IP multicast.
- IGMP is defined in RFC 1112.
- IGMP operates on a physical network (e.g., single Ethernet Segment).
- IGMP is used by multicast routers to keep track of membership in a multicast group.
- Support for:
 - Joining a multicast group
 - Query membership
 - Send membership reports

IGMP Protocol

- A host sends an **IGMP report** when it joins a multicast group (Note: multiple processes on a host can join. A report is sent only for the first process).
- No report is sent when a process leaves a group
- A multicast router regularly multicasts an **IGMP query** to all hosts (group address is set to zero).
- A host responds to an IGMP query with an **IGMP report**.
- Multicast router keeps a table on the multicast groups that have joined hosts. The router only forwards a packet, if there is a host still joined.
- Note: Router does not keep track which host is joined.


IGMP Packet Format

- IGMP messages are only 8 bytes long




- Type: 1 = sent by router, 2 = sent by host

IGMP Protocol



IGMP Protocol



IGMP general query

IGMP group address = 0

Destination IP address = **224.0.0.1**

Source IP address = router's IP address

IGMP group-specific query

IGMP group address = **group address**

Destination IP address = **group address**

Source IP address = router's IP address

IGMP membership report


IGMP group address = **group address**

Destination IP address = **group address**

Source IP address = host's IP address


Networks with multiple multicast routers

- Only one router responds to IGMP queries (**Querier**)
 - Router with smallest IP address becomes the querier on a network.
- One router forwards multicast packets to the network (**Forwarder**).




Multicast Routing Protocols

- **Goal:** Build a spanning tree between all members of a multicast group




Multicast routing as a graph problem

- **Problem:** Embed a tree such that all multicast group members are connected by the tree




Multicast routing as a graph problem

- **Problem:** Embed a tree such that all multicast group members are connected by the tree
- **Solution 1: Shortest Path Tree or source-based tree**
Build a tree that minimizes the path cost from the source to each receiver
 - *Good tree if there is a single sender*
 - *If there are multiple senders, need one tree per sender*
 - *Easy to compute*



Multicast routing as a graph problem

- **Problem:** Embed a tree such that all multicast group members are connected by the tree
- **Solution 2: Minimum-Cost Tree**
Build a tree that minimizes the total cost of the edges
 - *Good solution if there are multiple senders*
 - *Very expensive to compute (not practical for more than 30 nodes)*



Multicast routing in practice

- Routing Protocols implement one of two approaches:

1. Source Based Tree:

- Essentially implements Solution 1.
- Builds one shortest path tree for each sender
- Tree is built from receiver to the sender → reverse shortest path / reverse path forwarding

2. Core-based Tree:

- Build a single distribution tree that is shared by all senders
- Does not use Solution 2 (because it is too expensive)
- Selects one router as a “core” (also called “rendezvous point”)
- All receivers build a shortest path to the core → reverse shortest path / reverse path forwarding

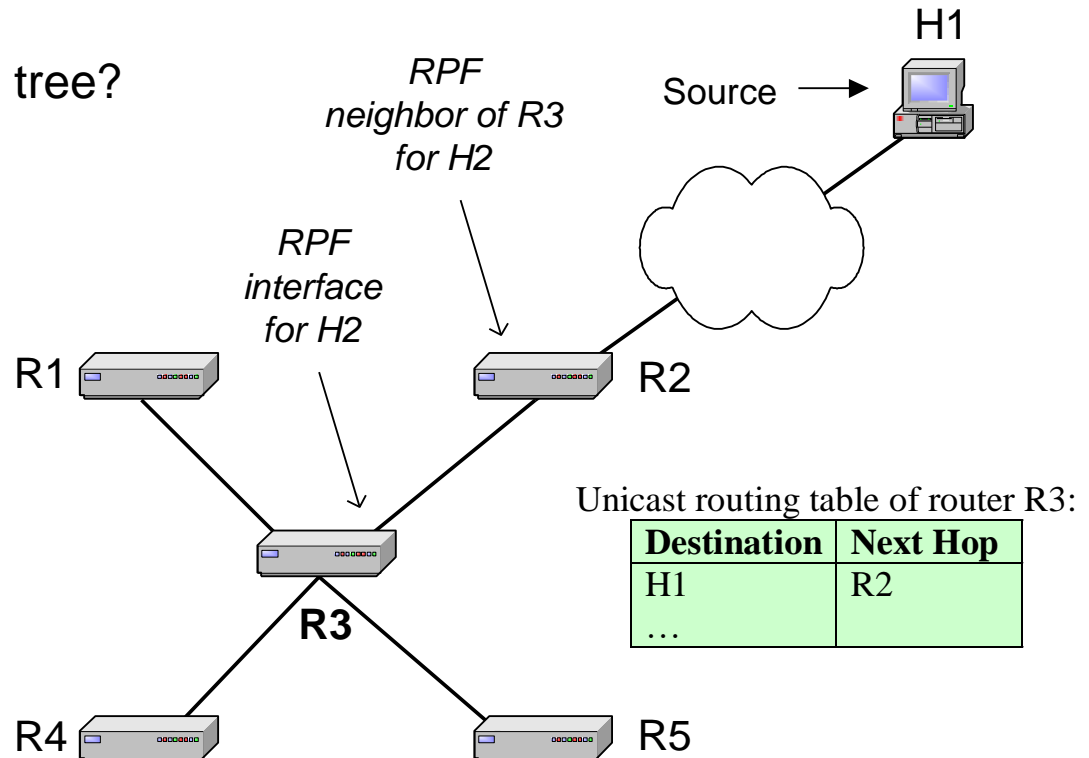
Reverse Path Forwarding (RPF)

- RPF builds a shortest path tree in a distributed fashion by taking advantage of the unicast routing tables.
- **Main idea:** Given the address of the root of the tree, a router selects as its upstream neighbor in the tree the router which is the next-hop neighbor for forwarding unicast packets to the root.
- How can this be used to build a tree?

1. RPF Forwarding:

Forward a packet only if it is received from an RPF neighbor

- ## 2. Set up multicast routing table in accordance from receiver to sender along the reverse shortest path tree



Multicast routing in practice

- Routing algorithms in practice implement one of two approaches:
 1. **Source Based Tree Tree:**
 - Establish a reverse path to the source
 2. **Core-based Tree:**
 - Establish a reverse path to the core router

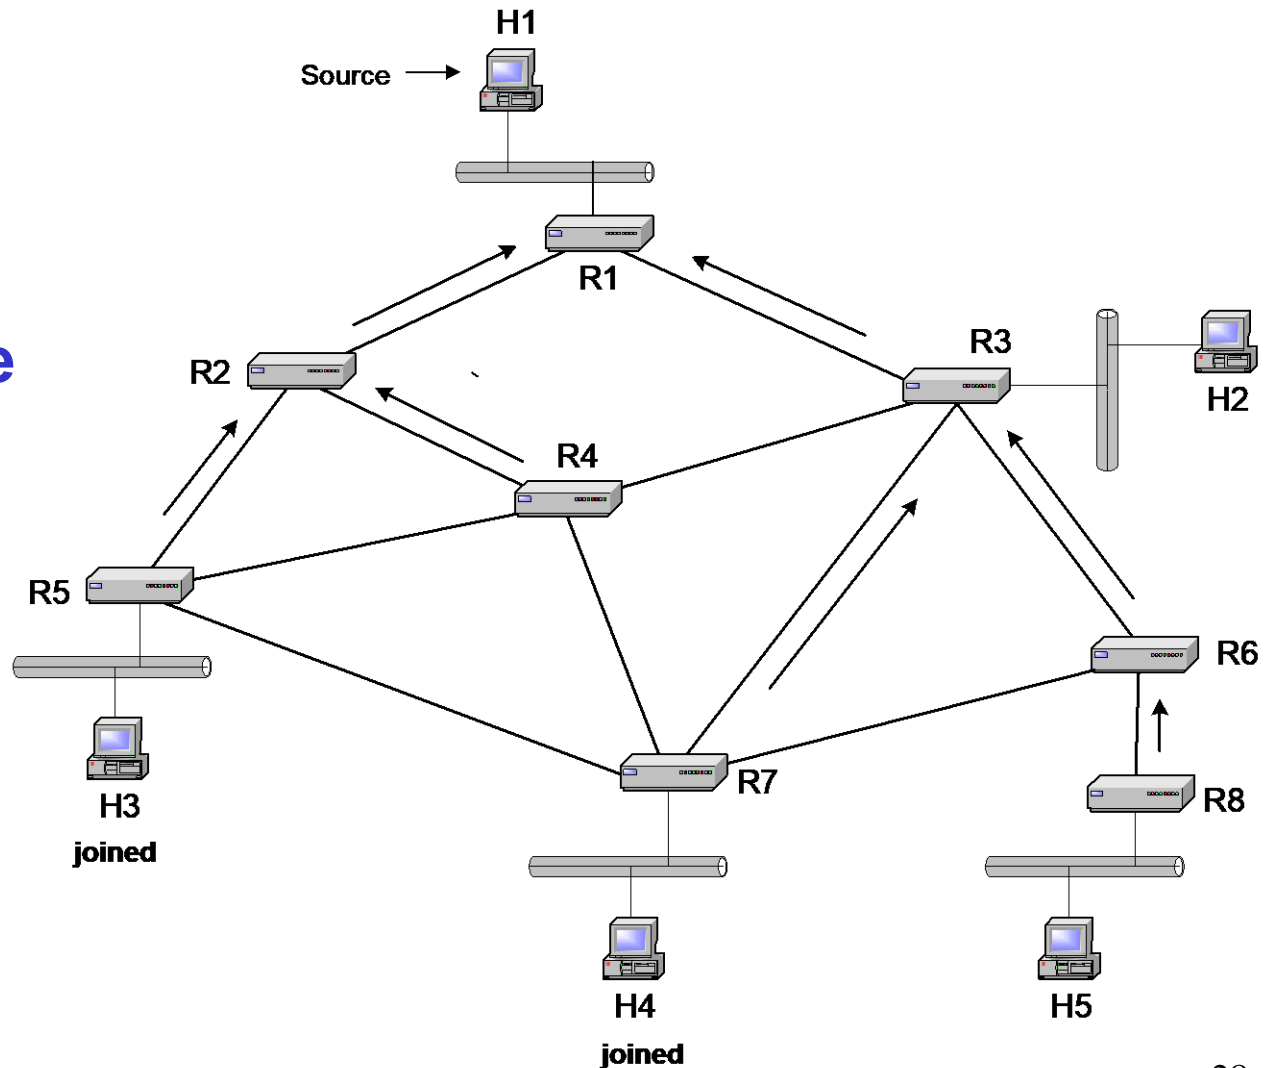
Multicast Routing table

- Routing table entries for source-based trees and for core-based trees are different
 - **Source-based tree**: (Source, Group) or (S, G) entry.
 - **Core-based tree**: (*, G) entry.

Source IP address	Multicast group	Incoming interface (RPF interface)	Outgoing interface list
S1	G1	I1	I2, I3
*	G2	I2	I1, I3

Building a source-based tree

- Set routing tables according to RPF forwarding
- **Flood-and-Prune**

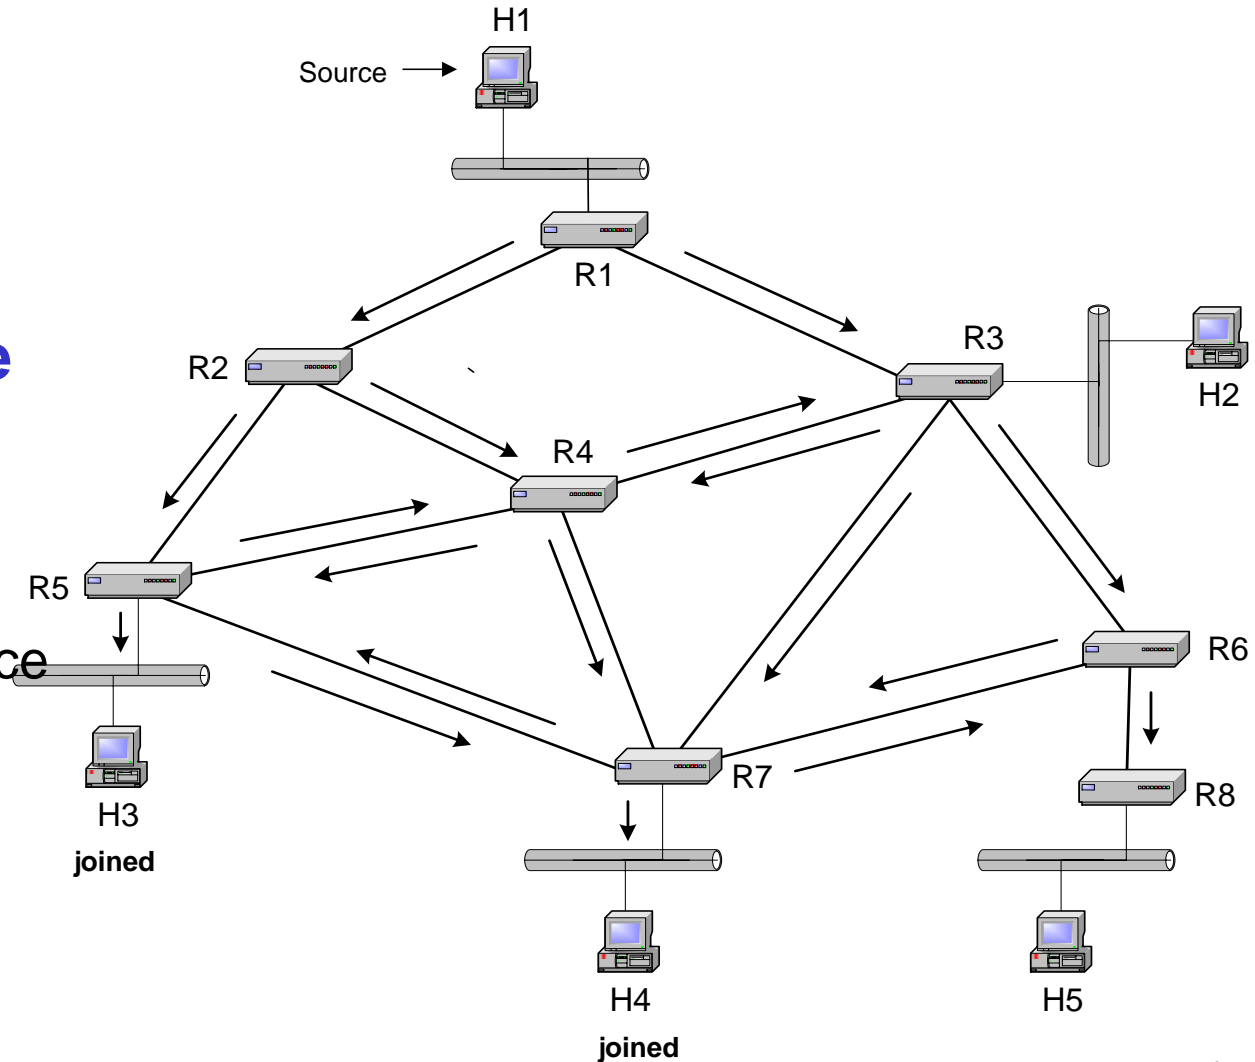


Building a source-based tree

- Set routing tables according to RPF forwarding
- **Flood-and-Prune**

Flood=

Forward packets that arrive on RPF interface on all non-RPF interfaces

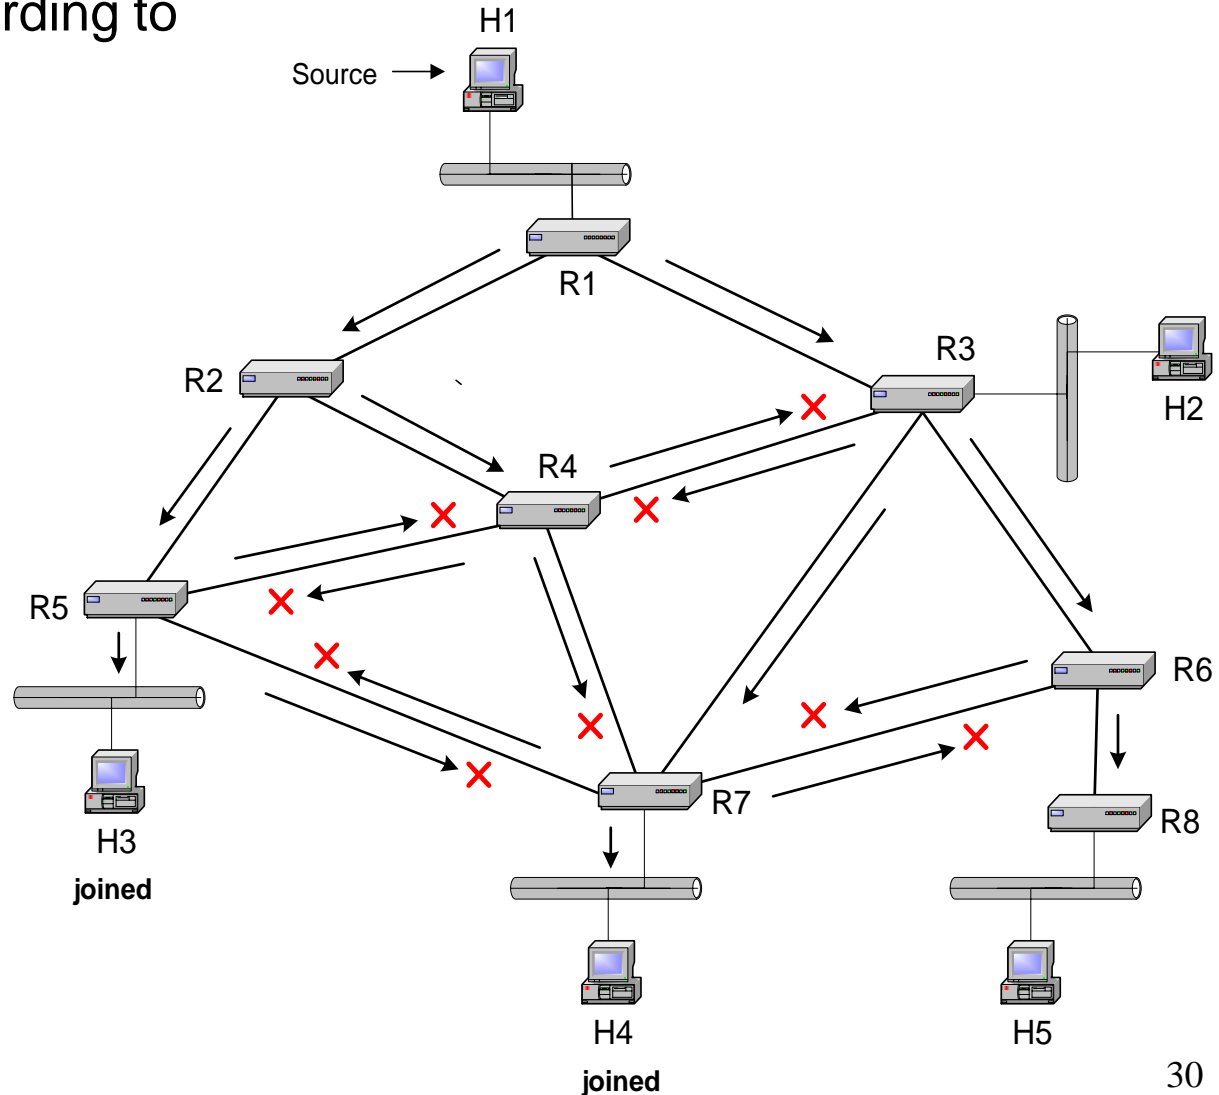


Building a source-based tree

- Set routing tables according to RPF forwarding
- **Flood-and-Prune**

Flood=
Forward packets
on all non-RPF interfaces

Receiver drops packets
not received on
RPF interface



Building a source-based tree


- Set routing tables according to RPF forwarding

- Flood-and-Prune**

Prune=

Send a prune message when a packet is received on a non-RPF interface or when there are no receivers downstream

Prune message disables routing table entry

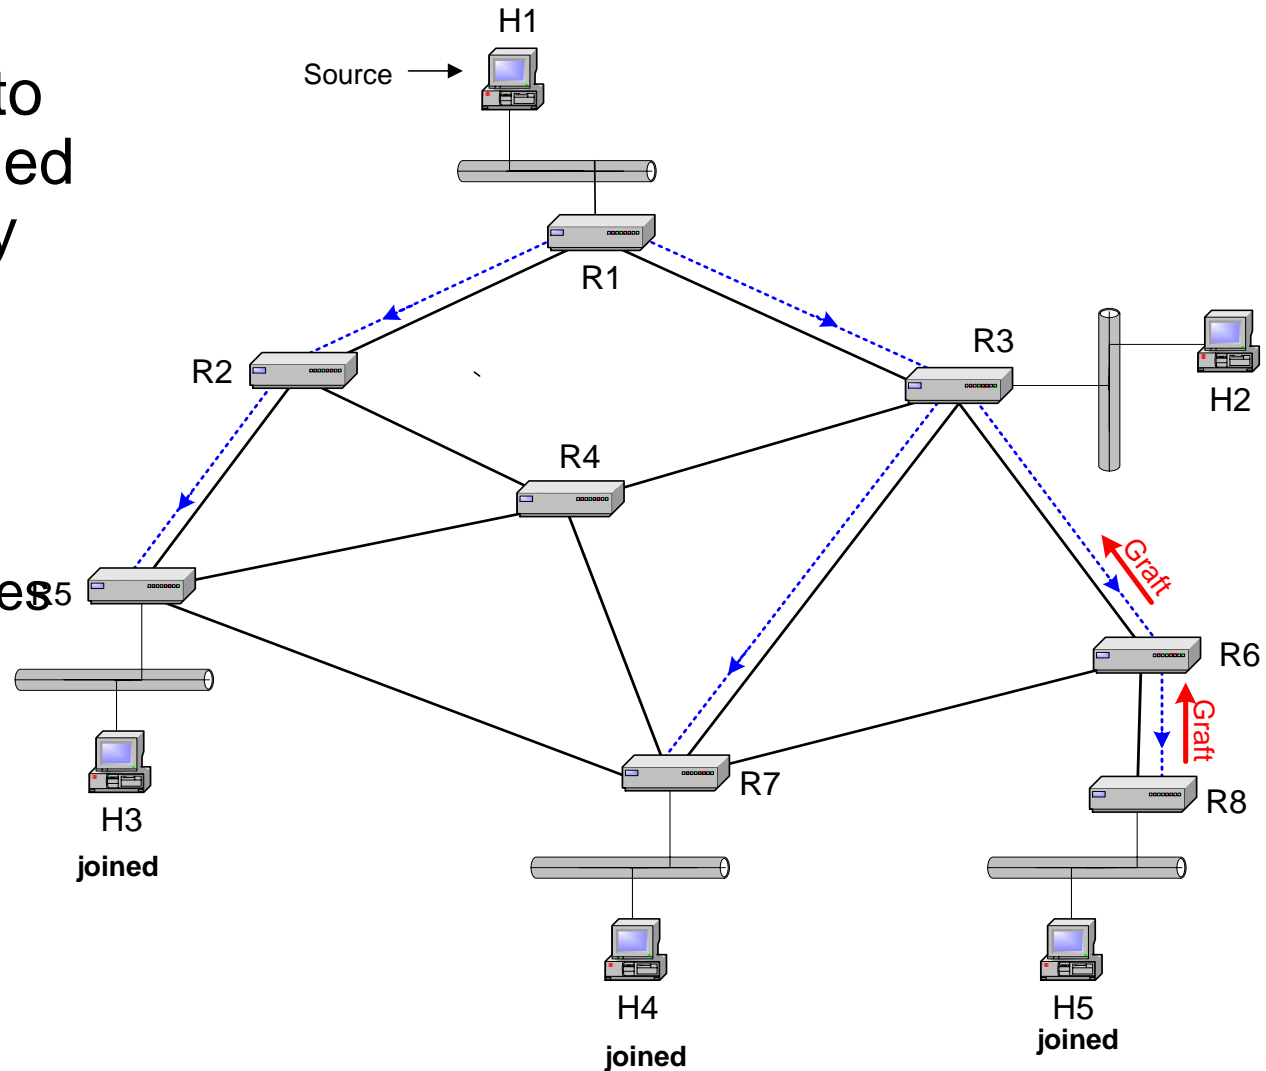


Pruning

- **Prune message** temporarily disables a routing table entry
 - **Effect:** Removes a link from the multicast tree
 - No multicast messages are sent on a pruned link
 - Prune message is sent in response to a multicast packet
 - *Question: Why is routing table only temporarily disabled?*
- Who sends prune messages?
 - A router with no group members in its local network and no connection to other routers
 - A router with no group members in its local network which has received a prune message on all non-RPF interfaces
 - A router with group members which has received a packet from a non-RPF neighbor

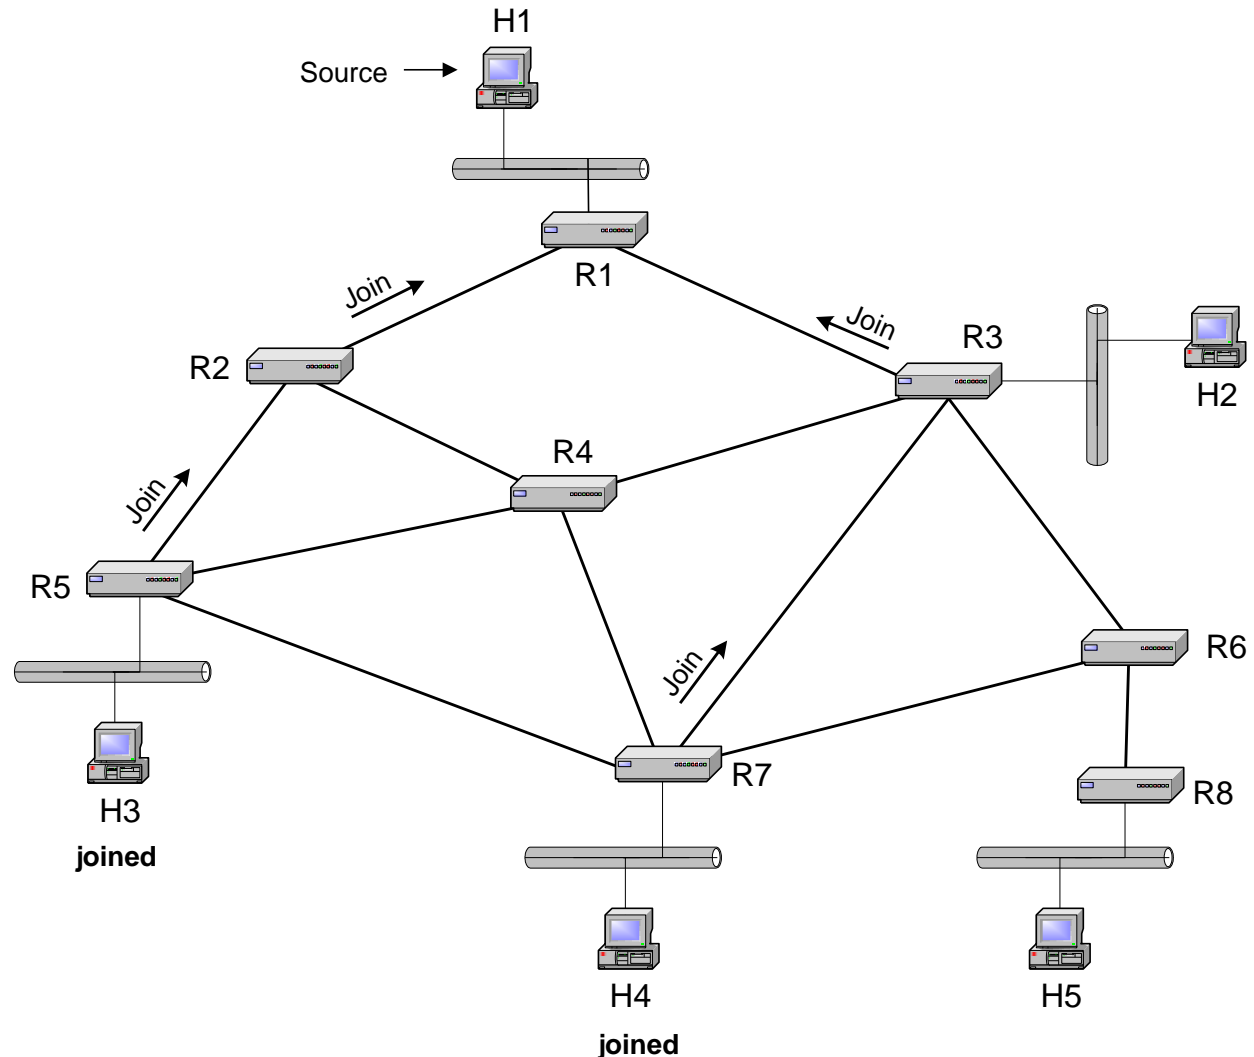
Building a source-based tree

- When a receiver joins, one needs to re-activate a pruned routing table entry
- Grafting**
Sending a Graft message disables prune, and re-activates routing table entry.




Alternative method for building a source-based tree

- This only works if the receiver knows the source
- **Explicit-Join**
 - Receiver sends a Join message to RPF neighbor
 - Join message creates (S,G) routing table entry
 - Join message is passed on




Building a core-based tree

- One route is the **core**
- Receiver sends a Join message to RPF neighbor with respect to core
- Join message creates (*, G) routing table entry



Building a core-based tree

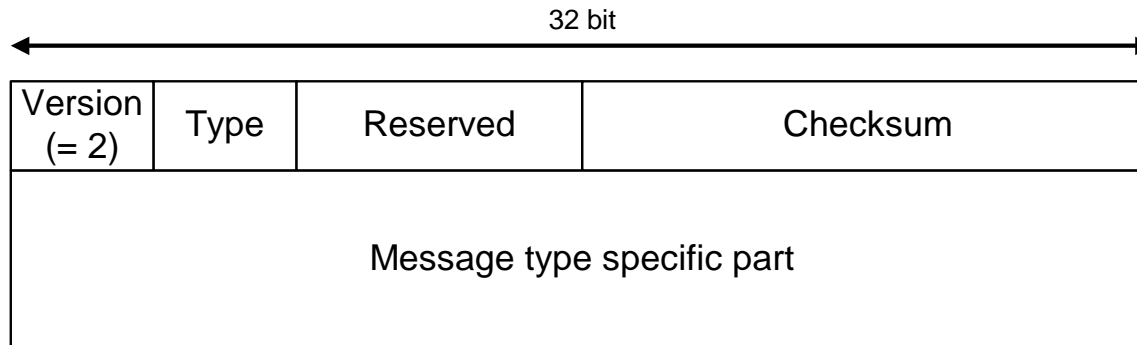
- Source sends data to the core
- Core forwards data according to routing table entry



Multicast routing protocols in the Internet

- **Distance Vector Multicast Routing Protocol (DVMRP):**
 - First multicast routing protocol
 - Implements flood-and-prune
- **Multicast Open Shortest Path First (MOSPF):**
 - Multicast extensions to OSPF. Each router calculates a shortest-path tree based on link state database
 - Not widely used
- **Core Based Tree (CBT):**
 - First core-based tree routing protocol
- **Protocol Independent Multicast (PIM):**[\[1\]](#)
 - Runs in two modes: PIM Dense Mode (PIM-DM) and PIM Sparse Mode (PIM-SM).
 - PIM-DM builds source-based trees using flood-and-prune
 - PIM-SM builds core-based trees as well as source-based trees with explicit joins. [\[1\] RFC2362](#)

PIM Messages (PIM version 2)

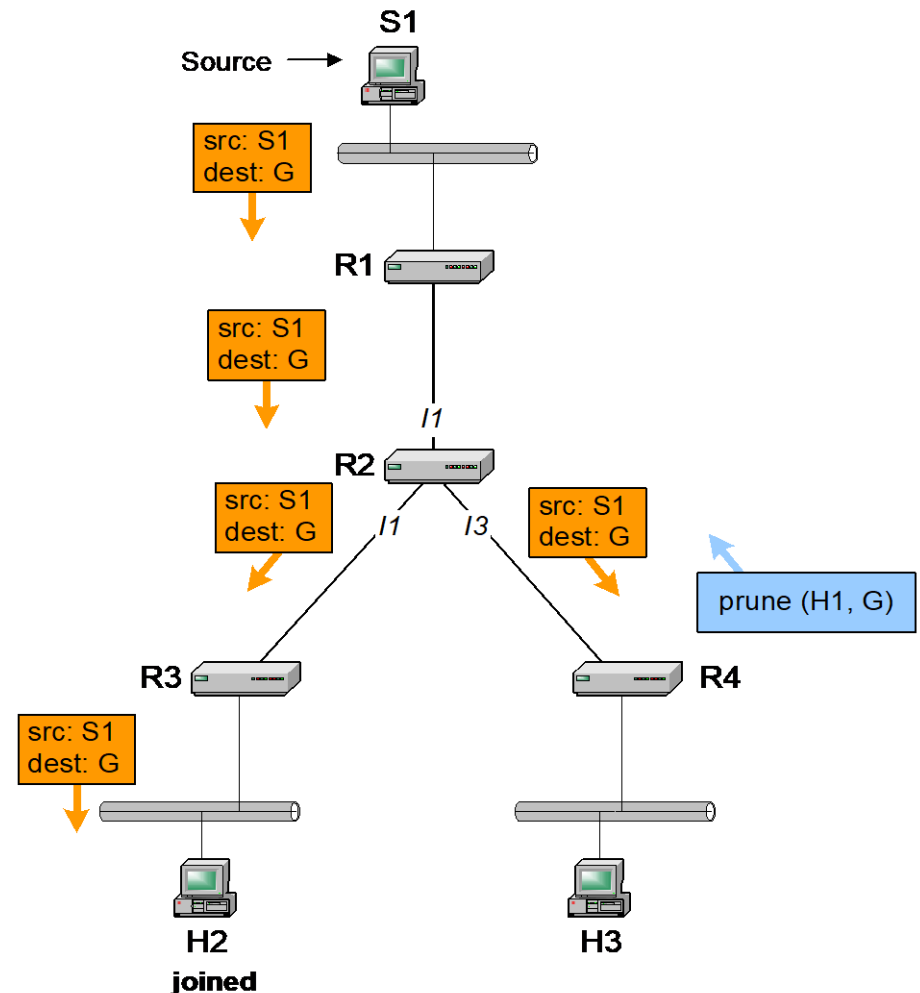


PIM-DM messages	Type	PIM-DM	PIM-SM
Hello	0	✓	✓
Register	1		✓
Register-Stop	2		✓
Join/Prune	3	✓	✓
Bootstrap	4		✓
Assert	5	✓	✓
Graft	6	✓	
Graft-Ack	7	✓	
Candidate-RP- Advertisement	8		✓

- Encapsulated in IP datagrams with protocol number 103.
- PIM messages can be sent as unicast or multicast packet
- 224.0.0.13 is reserved as the *ALL-PIM-Routers* group


PIM-DM: PIM Dense Mode

- PIM-DM implements flood-and-prune
- Orange packet: Multicast packet (=Data)
- Blue packet: PIM message



PIM-SM: PIM Sparse Mode


- Core is called rendezvous-point (**RP**)
- Receivers know RP (statically configured or dynamically elected)
- When receiver joins, a Join message is sent to RP on RPF.



(a) PIM-SM: H2 joins

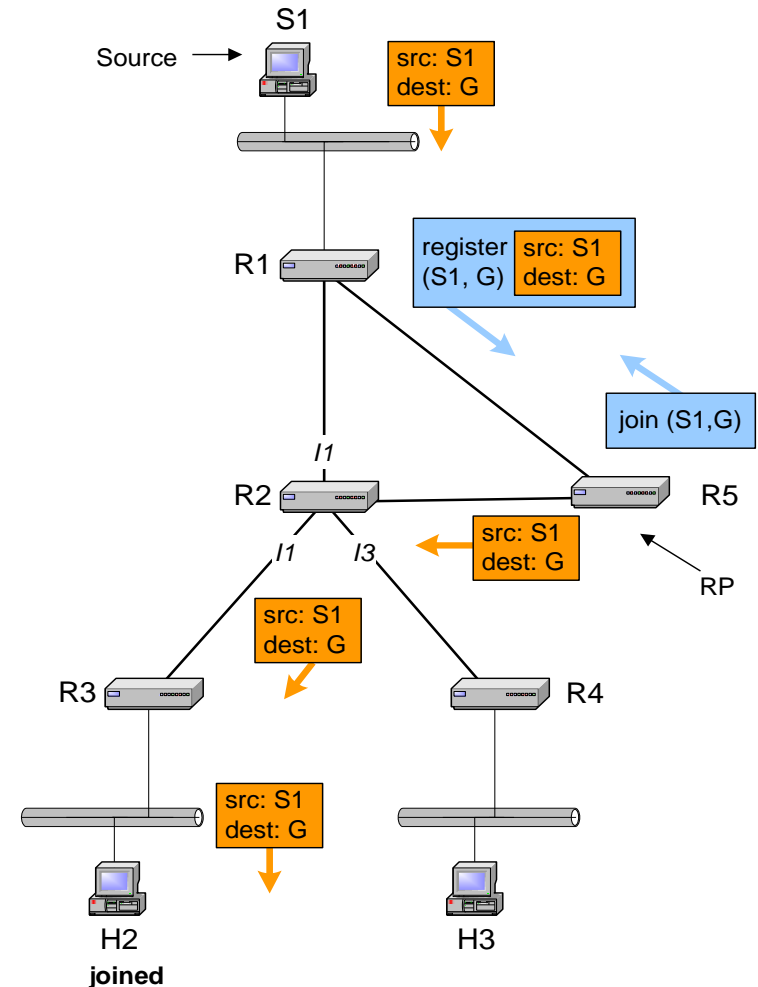
PIM-SM: PIM Sparse Mode

- Host H3 joins:
Join message is only forwarded until the first router that is part of the core-based tree.



PIM-SM: Data transmission

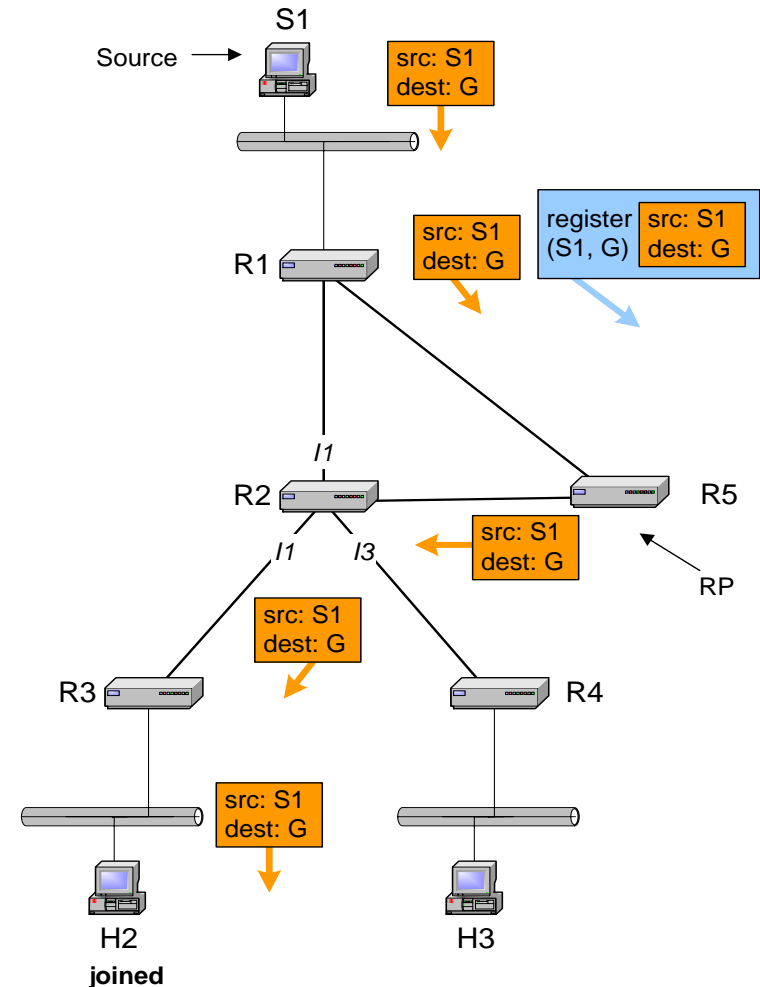
- Source sends multicast packet to RP
- Packet is attached to an RP Register message
- When packet reaches RP, it is forwarded in the tree
- Also: RP sends a Join message on reverse path to S1



(a) PIM-SM: Register message to RP

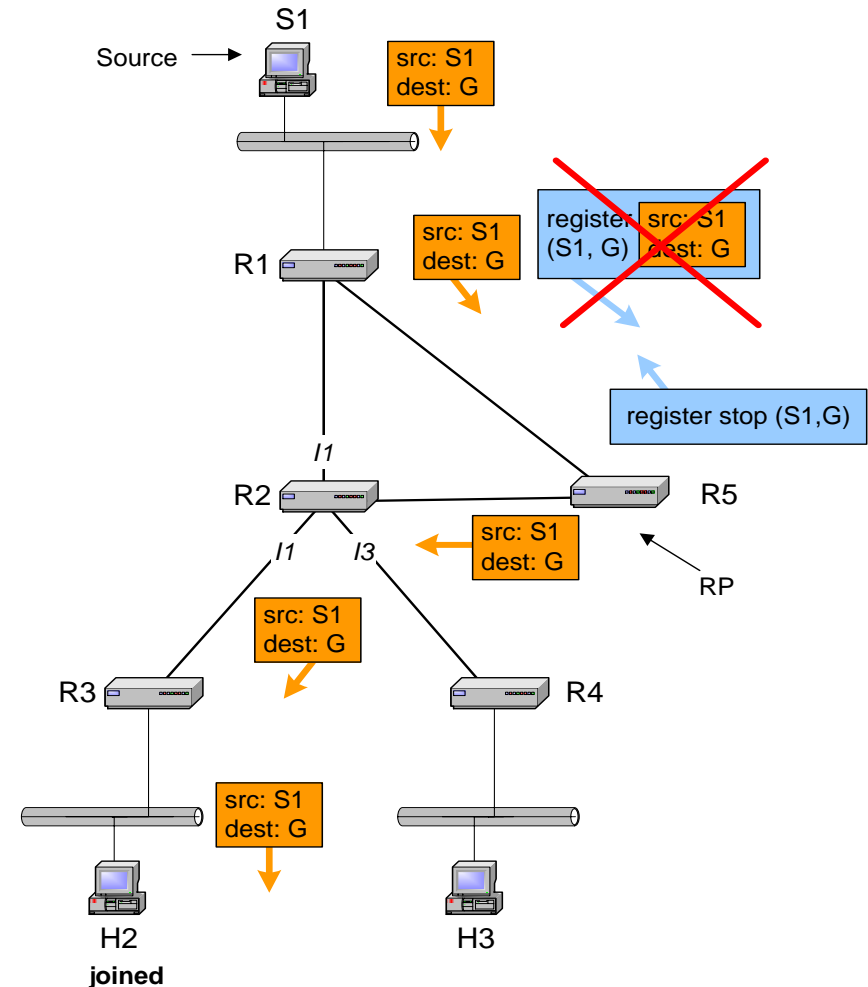
PIM-SM: Data transmission

- When Join messages reaches R1, it sends a native multicast packet to the RP (in addition to the packet attached to the register message)




PIM-SM: Data transmission

- When RP receives native multicast packet it sends a register stop message to R1. This message stops the transmission of register messages from R1.



PIM-SM: Switching to source-based tree


- When data to receivers exceeds a threshold, routers switch to a source-based tree
- This is done by sending an explicit join message to the source
- There may be duplicate packets being sent for some time



(a) PIM-SM: R3 switches to a SPT

PIM-SM: Switching to source-based tree

- When data arrives from source (as opposed to RP), a Prune message is sent to the RPT
- Now: data is forwarded only along the shortest-path tree



(b) PIM-SM: Data follows a SPT